JOB

Interchained Job Scout

Generated Saturday, April 18, 2026 · 130 leads queued

Founding Software Engineer

100% fit

Wavelength · ? · On-site · ycombinator

ROLE

#

Founding Software Engineer at Wavelength (W24)

$120K - $160K • 0.20% - 0.70%

Build Magical Customer Experiences

San Francisco, CA, US Full-time US citizen/visa only Any (new grads ok)

Apply now

About the role

What you'll do

You'll be responsible for shipping products from 0-1 and their maintenance once they're in production. We're a customer obsessed team and we also expect you to talk to our customers, iterate on what you push, and continuously improve. Most of all, you will be the sole architect, designer, and engineer on multiple systems across our stack

Technical Must-Have's

- High Level Proficiency in Go, Typescript, Python, Java, C/C++

- Designing / Working on Distributed Systems

Brownie Points (not required, but encouraged)

- Experience / Proficiency with Rust,

- Frontend Experience (React, Next.Js, similar)

What a Good Fit looks like

- Enjoy solving open-ended, challenging problems

- Value good design and UX (an eye for it)

- Thrive with owning end-to-end product development

- Able to easily communicate complex ideas simply

- Extremely Proactive (about everything)

What we Care about

- Customer Obsession: Our goal is to make every customer love what we build, whatever it takes.

- Beautiful Design: Software should be beautiful, it's in the details

- Accountability: Taking things end-to-end, owning up to when we make mistakes, and fully trusting each other along the way

- HighBar: We demand a high bar for ourselves, we expect you to have one to…

https://www.workatastartup.com/jobs/80203

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Founding Engineer role — passing this along if it lands in the wrong spot

---

Hey —

Wavelength's 0-to-1 customer experience work caught my attention — specifically the idea of engineers staying close to users while owning full product cycles. That's the mode I've been in.

Most relevant thing I've shipped: a multi-provider AI platform with routing, governance, and enterprise control surfaces across 11+ model providers — built from scratch, maintained in production, iterated directly from user feedback.

That loop — build, talk to users, improve — is exactly how I work.

If you're not the right person to send this to, no pressure — happy if you can point me toward whoever owns engineering hiring. Otherwise, would love 20 mins.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:80203 · captured: 2026-04-19 · fit_reason: —

Founding Engineer (Fullstack)

Onlook · ? · On-site · ycombinator

100% fit

ROLE

#

Founding Engineer (Fullstack) at Onlook (W25)

$130K - $200K • 1.00% - 4.00%

Cursor for Designers

San Francisco, CA, US / Remote (US; CA) Full-time US citizen/visa only Any (new grads ok)

Apply now

About Onlook

Onlook is on a mission to bring more people together in code. We built a visual editor that anyone can use to contribute code to real codebases – or, just create a website super fast.

Since launching, Onlook hit #1 on Hacker News, was the #1 trending repo on GitHub—above DeepSeek, Anthropic—and has earned over 19,000 GitHub stars.

If you're looking for an opportunity to be a part of a next-generation design-led company, Onlook might be the right place for you.

About the role

Skills: Next.js, Tailwind, Node.js, React, TypeScript, PostgreSQL

Since launching, Onlook hit #1 on Hacker News, was the #1 trending repo on GitHub – above DeepSeek + Anthropic – and has earned over 22,000 GitHub stars. We’re looking to bring on Onlook’s first Founding Engineer. You’re a full-stack engineer who is ready to build a generational design-led company.

This role requires autonomy - you’ll be setting standards for one of the fastest-growing open source projects backed by YC ever. You’ll help design and build an uncompromising visual IDE loved by tens of thousands of designers and engineers around the world.

###

Why Onlook exists

Software creation is still broken at its core. Design and engineering operate in parallel, constrained by single-player tools that limit explorat…

https://www.workatastartup.com/jobs/74305

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Founding Engineer role — passing this along if it lands wrong

Hey —

Watching Onlook climb to #1 on HN and top GitHub trending — above DeepSeek — while solving the design-to-code gap in real codebases is genuinely interesting work.

I'm Mark. I recently architected a multi-provider AI platform with routing, governance, and enterprise control surfaces across 11+ model providers — fullstack, TypeScript/Node-heavy, with real API surface design. Feels

adjacent to what your team is scaling.

I'd love 20 minutes with whoever owns engineering hiring — or if that's not you, happy for you to forward this along.

GitHub: github.com/aiassistsecure
Portfolio: AiAssist.net

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:74305 · captured: 2026-04-19 · fit_reason: —

Founding Engineer, Full Stack

ProhostAI · ? · On-site · ycombinator

100% fit

ROLE

#

Founding Engineer, Full Stack at ProhostAI (S24)

$100K - $180K • 0.20% - 2.00%

AI Property Manager for Airbnb & Short-Term Rental Hosts

San Francisco, CA, US / Remote (US) Full-time Will sponsor Any (new grads ok)

Apply now

About ProhostAI

We are a YC S24 company on a mission to make every Airbnb host a 5-star host by automating and improving their operations with the help of AI. We are building an AI-powered operating system for hosts on Airbnb and other short-term rental platforms.

Our founding team consists of highly experienced technical leaders from Airbnb and Dropbox. We are looking to expand our small but mighty team.

About the role

Company Description

ProhostAI (YC S24) building an AI-powered operating system for Airbnb & short-term rental hosts. Our founding team consists of data scientists and engineers from Airbnb, Dropbox, Twitter and Harvard. We are backed by Y Combinator, Pear VC, Pioneer Fund, Data Tech Fund and others. Our founding team comes from Airbnb and Dropbox engineers and data scientists who are also hosts with a 4.99+ rating across thousands of stays.

Role Description

This is a full-time hybrid role based out of San Francisco. As a Full Stack Product Engineer, you will be responsible for end to end back-end and front-end web development. You will also work on AI/LLM powered features throughout our product. We are looking for a highly driven, product minded engineer who is excited to ship novel & delightful AI-powered user experiences to…

https://www.workatastartup.com/jobs/68411

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Founding Engineer role — passing this along if it lands wrong

Hey —

ProhostAI's angle of building a full OS for STR hosts — not just a messaging bot — caught my attention. Automating the operational layer across guest comms, listing optimization, and host workflows is the right scope for this problem.

I'm a full-stack AI systems engineer with 4 years building multi-provider AI platforms — most recently architecting a BYOK routing system with governance and control surfaces across 11+ model providers. That kind of infrastructure maps pretty directly to what you're building.

If you're involved in hiring, happy to find 20 mins. If not, no pressure — would appreciate a point in the right direction.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:68411 · captured: 2026-04-19 · fit_reason: —

Founding Engineer

Orange Slice · ? · On-site · ycombinator

100% fit

ROLE

#

Founding Engineer at Orange Slice (S25)

$150K - $250K • 0.50% - 5.00%

Agentic sales enrichment spreadsheet

San Francisco, CA, US / Remote (US) Full-time US citizen/visa only Any (new grads ok)

Apply now

About the role

Skills: TypeScript

##

Founding Engineer — Orange Slice (YC)

###

About Orange Slice

Orange Slice helps revenue teams prospect, enrich, qualify, and automate GTM work with AI . Customers use it to turn messy intent and company context into actionable lists, enriched records, and workflows —from firmographics and contact verification to web research, social signals, and outreach drafting. The product spans a polished web experience (including spreadsheet-like execution and an AI agent) and tooling for builders who want the same capabilities in code or in Claude Code via a single install path.

We sit at the intersection of B2B data , LLMs , and workflow automation —where reliability, cost control, and great UX matter as much as model quality.

###

The role

We’re hiring a Founding Engineer to own major slices of the product and infrastructure alongside the founders. This is a high-ownership, high-trust role: you’ll ship user-facing features, harden core systems, and help define how we scale—not just close tickets.

You might work on:

-

Core product : Next.js App Router surfaces (dashboard research tools, spreadsheet workflows, agent/chat UX), performance, and end-to-end flows from signup to paid usage.

-

Data + AI execution : orchestrating long-r…

https://www.workatastartup.com/jobs/85602

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Founding Engineer interest — Orange Slice

Hey,

What you're building at Orange Slice — turning raw intent signals into enriched, actionable GTM workflows through an agent layer — is exactly the problem worth solving right now.

I'm Mark. I recently architected a multi-provider AI platform with tool routing, memory-aware interaction patterns, and developer-facing API surfaces across 11+ model providers — pretty close to the agent-plus-spreadsheet surface Orange Slice is building on.

TypeScript is my primary language. Portfolio and GitHub are at AiAssist.net and github.com/aiassistsecure if useful context.

If you're not the right person to loop in, happy for you to forward this along — or just point me toward whoever owns engineering hiring.

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:85602 · captured: 2026-04-19 · fit_reason: —

Director of Platform Engineering

75% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States / Remote (US) · Remote ' · ycombinator

Ø=Ü° $253K - $352K

ROLE

Director of Platform Engineering at Checkr | Y Combinator Checkr People infrastructure for the future of work Director of Platform Engineering $253K - $352K • Denver, Colorado, United States; San Francisco, California, United States / Remote (US) Job type Full-time Role Engineering, Full stack Experience 11+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the role As the Director of Platform Engineering at Checkr, you will be responsible for driving the organization's overall platform and infrastructure strategy. To achieve this, you will need to have a deep understanding of platform engineering and how to drive a strategic roadmap of foundational capabilities and, most importantly, its adoption by the other development teams. What you'll do Guide the technical vision and strategy for the best-in-class highly scal…

https://www.ycombinator.com/companies/checkr/jobs/KWiLgmX-director-of-platform-engineering

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Platform Engineering – Mark Allen Evans

Hi,

Checkr's push to own compliance and risk infrastructure at the scale of DoorDash and Coinbase — while landing Forbes Cloud 100 this year — caught my attention in a real way.

I'm Mark Allen Evans. I recently architected a multi-provider AI platform with routing, governance, and enterprise control surfaces across 11+ model providers — the kind of platform ownership this Director role seems to center on.

I'd welcome 20 minutes with whoever owns this search — or happy to send more context your way if that's easier.

If you're not close to hiring, no pressure — feel free to pass this along.

Mark Allen Evans
https://AiAssist.net | https://github.com/aiassistsecure

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92267 · captured: 2026-04-18 · fit_reason: keyword + remote match

Cloud Engineer (Remote)

Joveo Ai · United Kingdom (Remote) · Remote ' · linkedin_jobs

75% fit

ROLE

About Joveo:Every company says they're "AI-first." We actually are. Joveo's recruitment advertising platform processes millions of hiring decisions through machine learning, real-time bidding, and predictive analytics — helping the world's largest employers find the right people, faster and fairer. But we're not done. Not even close.
Role: Cloud EngineerLocation: Remote
Role Overview:We are hiring a Cloud Engineer to design, deploy, and manage cloud-native infrastructure that powers Joveo's global platform. You will build resilient, scalable, and cost-optimized cloud systems — ensuring our infrastructure grows with the business without compromising performance or reliability.
Key Responsibilities:- Architect and deploy cloud infrastructure across AWS, GCP, or Azure- Implement infrastructure-as-code using Terraform or similar tooling- Optimize cloud resource utilization and manage cost governance- Design networking, security groups, IAM policies, and access controls- Support engineering teams with cloud migrations and modernization projects- Build runbooks and disaster recovery procedures for critical systems
Required Skills & Qualifications:- Deep hands-on experience with at least one major cloud platform- Proficiency in Terraform, Pulumi, or AWS CDK for infrastructure provisioning- Knowledge of networking fundamentals — VPCs, subnets, DNS, load balancers- Experience with cloud security best practices and compliance frameworks- Familiarity with Kubernetes, ECS, or serverless …

https://www.linkedin.com/jobs/view/4403433592

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Cloud Engineer role at Joveo — Mark Allen Evans

Hi,

Joveo's real-time bidding and ML pipeline running at global hiring scale is exactly the kind of infrastructure problem I find most interesting right now.

I'm a Cloud/AI Systems Engineer with 4 years building production-grade platforms. Most relevant here: I designed secure runtime separation between customer-facing APIs and isolated execution services using signed internal requests, allowlisted access, and full audit trails — the kind of boundary enforcement that matters when you're processing millions of decisions at scale.

GitHub: github.com/aiassistsecure | Portfolio: AiAssist.net

If you're not close to hiring for this, happy for you to forward it along — or I'd love 20 mins with whoever owns it.

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4403433592 · captured: 2026-04-18 · fit_reason: keyword + remote match

DevOps Engineer, Trilogy (Remote) - $100,000/year USD

75% fit

Crossover · London, England, United Kingdom (Remote) · Remote ' · linkedin_jobs

ROLE

You're the engineer who maintains uptime for 50+ SaaS products when others are still guessing. We need DevOps engineers who can step into unfamiliar AWS environments, bring order to instability, and drive availability beyond 99.9% through disciplined monitoring, automation, and root cause analysis. You'll break down complex projects into executable one-day increments, deliver production-ready Python or JavaScript, and leverage AI as your assistant.

Most organizations talk about "cloud-native" while manually tending infrastructure. We're systematizing reliability across dozens of acquired products where original teams have departed and documentation is incomplete. The challenge: you'll apply agents and contemporary tooling to understand new systems 5–10x faster, document your findings, and automate solutions so repeat failures are eliminated. Rather than judging you on certifications and buzzwords, we'll observe how you troubleshoot under pressure, author a genuine 5-Whys that identifies one preventable root cause, and construct automations that endure in production.

This is not a tier-2 "follow the runbook" position. Here, you author the runbooks, architect the deployment path from development through staging to gradual rollout with soak periods and rollback criteria, and implement monitoring that detects corner cases. You reject risky changes before they reach production. You distinguish infrastructure failures you own from application bugs Engineering owns, and you route …

https://www.linkedin.com/jobs/view/4395939580

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: DevOps Engineer role – worth a quick look?

Hi,

Crossover's model of maintaining 99.9%+ uptime across 50+ SaaS products simultaneously — without the luxury of familiarity — is exactly the kind of constraint that separates disciplined engineers from the rest.

I've been doing something adjacent: at AiAssist Secure, I designed secure runtime separation between customer-facing APIs and isolated execution services using signed internal requests, allowlisted access, and auditable actions across a multi-provider AI platform. Stability under complexity is the job.

I write production Python and JavaScript, work natively in AWS, and use AI tooling daily — not as a novelty.

If you're not the right person to send this to, no pressure — happy if you can point me in the right direction. Otherwise, I'd welcome 20 mins with whoever owns this.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4395939580 · captured: 2026-04-18 · fit_reason: keyword + remote match

GCP/AWS Platform Engineer - start-up experience

75% fit

Onsera Health · Greater London, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

The ChallengeCardiometabolic conditions will impact 80-90% of people throughout their lifetimes, representing the leading cause of death globally and an important risk factor for a variety of neurodegenerative diseases and cancer. Sustainably addressing these diseases requires breakthrough therapeutics, AI/ML innovation, and transformative business models that translate clinical outcomes into economic value.The OpportunityOwn the cloud platform that powers Onsera Health's breakthrough healthcare AI. We are backed by Population Health Partners (PHP), the proven venture platform behind Metsera and Corsera Health, and, together, we're building the future of population health management.As our Platform Engineer, you will build and operate production-grade GCP infrastructure that enables fast and secure product delivery across all our global engineering and AI teams. You will be in the central seat across all our deliveries including data, ML, LLM, agentic AI, and applications. ResponsibilitiesOwn GCP infrastructure – Operate and evolve production and non-production environments; manage cost visibility, quotas, and alerts; maintain clean environment separation with appropriate access controlsImplement Infrastructure as Code – Build and maintain Terraform across projects, networks, IAM, and servicesEnforce security and compliance – Implement controls aligned with healthcare requirements (HIPAA, SOC 2, ISO 27001); configure IAM, service accounts, workload identity, network security,…

https://www.linkedin.com/jobs/view/4381746616

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Platform Engineer role at Onsera Health

Hi,

Onsera's dual-cloud build across GCP and AWS caught my attention — that kind of hyperscaler-agnostic architecture in health-tech is genuinely interesting engineering, and it signals the right kind of pragmatism early on.

I'm Mark — AI Systems Engineer with 4 years building platform-level infrastructure. One thing that maps directly: I designed secure runtime separation between customer-facing APIs and isolated execution services using signed internal requests, allowlisted access, and full audit trails — exactly the kind of control surface a regulated health platform needs.

If you're not involved in hiring, no worries — happy if you'd point me to whoever owns it. Otherwise, I'd welcome 20 minutes.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4381746616 · captured: 2026-04-18 · fit_reason: keyword + remote match

Cloud Platform Engineer - AI Trainer

DataAnnotation · Canada (Remote) · Remote ' · linkedin_jobs

75% fit

ROLE

Join the DataAnnotation team and contribute to developing cutting-edge AI systems, while enjoying the flexibility of remote work and setting your own schedule.

We are looking for proficient programmers to help advance AI development. As a member of DataAnnotation's coding team, you'll be part of a growing community of over 100,000 professionals — including front-end, back-end, full-stack, machine learning, and other engineers — who are driving real-world impact in AI development.

Our platform offers an engaging blend of flexibility and challenge: you'll work closely with state-of the art AI models to take on programming tasks that include creating and solving challenging coding problems, building beautiful apps with rich functionality, and synthesizing insights through data analysis and visualization. Your work directly contributes to refining intelligent systems that learn, adapt, and evolve. Some team members fit this work alongside a full-time role, while others treat it as their primary focus, choosing projects and schedules that align with their availability and goals.

To get started, once you sign up for an account, you'll take a short assessment (this serves as our version of an interview). If you pass that assessment, you'll receive an email confirmation, and paid work will become available to you through our platform.

Benefits

Fully remote: work from anywhere in the US, Canada, UK, Ireland, Australia, and New Zealand.Flexible schedule: choose which projects you …

https://www.linkedin.com/jobs/view/4397389632

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Cloud Platform Engineer – AI Trainer | Mark Allen Evans

Hi,

DataAnnotation's work sits at a layer most people overlook — the infrastructure that actually shapes how models learn. That's the part I find most interesting right now.

I'm a platform engineer with four years building AI systems end-to-end. One thing that maps directly here: I architected a BYOK multi-provider AI platform spanning 11+ model providers with routing, governance, and enterprise control surfaces — exactly the kind of structured, evaluable AI work your team runs at scale.

If you're not the right person to send this to, no pressure — happy if you can point me in the right direction. Otherwise, I'd welcome 20 minutes with whoever owns this.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4397389632 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI-Native Software Engineer - CDI Paris - Theodo

70% fit

Theodo · Greater Paris Metropolitan Region (Hybrid) · On-site · linkedin_jobs

ROLE

L'histoire de Theodo et son succès
Le groupe Theodo accompagne depuis 2009 les entreprises innovantes dans la conception, le développement et le déploiement de produits digitaux ingénieux, en tirant parti du meilleur de la technologie et de l'approche Lean.

Notre approche unique est fondée sur 3 principes :

nous avons développé des expertises reconnues dans les technologies clés de la conception et du développement produit; quelle que soit l'étape du projet, notre priorité à chaque décision est la création de valeur pour nos clients; nous recrutons les meilleurs talents digitaux et cultivons leur ingéniosité grâce à l'approche lean

Theodo connait une croissance exceptionnelle depuis 15 ans : nos équipes rassemblent plus de 700 Theodoers, principalement des Software Engineers, passionnés de technologie et d'amélioration continue. Avec la croissance, le groupe s'est structuré en entités spécialisées par technologies (Theodo Apps, Theodo Cloud, Theodo Data & AI) ou par industries (Fintech, Govtech, Healthtech, Consumer Products Goods). En 2025, le groupe Theodo génère 100M€ de CA.

Ton équipe :
Tu rejoindras une communauté tech de plus de 120 Theodoers : software engineers, tech leads, engineering managers et architectes, avec lesquels tu pourras échanger, apprendre et progresser au quotidien.

Chez Theodo, les techs travaillent en équipes projet, constituées en fonction des besoins de chaque client. Tu feras partie d'une équipe de Theodoers à taille humaine (4 à 5 perso…

https://www.linkedin.com/jobs/view/4375103504

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: AI-Native Engineer role @ Theodo — passing this along if it lands wrong

Hi,

Theodo's move to make AI a core delivery competency — not a bolt-on — caught my attention. That's a real architectural shift, not a trend play.

I'm Mark. I recently architected a BYOK multi-provider AI platform spanning 11+ model providers with routing, governance, and enterprise control surfaces — built to ship, not demo.

That work maps pretty directly to what your AI-native teams are delivering for clients.

If you're not the right person to send this to, happy for you to forward it along — or just point me in the right direction.

Either way, would love 20 mins with whoever owns this.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4375103504 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI Infrastructure Engineer

Intercom · London, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

70% fit

ROLE

Intercom is the AI Customer Service company on a mission to help businesses provide incredible customer experiences.

Our AI agent Fin, the most advanced customer service AI agent on the market, lets businesses deliver always-on, impeccable customer service and ultimately transform their customer experiences for the better. Fin can also be combined with our Helpdesk to become a complete solution called the Intercom Customer Service Suite, which provides AI enhanced support for the more complex or high touch queries that require a human agent.

Founded in 2011 and trusted by nearly 30,000 global businesses, Intercom is setting the new standard for customer service. Driven by our core values, we push boundaries, build with speed and intensity, and consistently deliver incredible value to our customers.

What's the opportunity?

We’re looking for Senior+ AI Infrastructure Engineers to build the systems that train and serve Intercom’s next generation of AI products.

Intercom is an AI company that builds from the GPU all the way up to a user agent that resolves millions of customer service queries a month.

You’ll join a small, highly technical team working at the cutting edge of modern AI infrastructure. The AI Infra team built the training pipelines and runs the inference for custom models like Fin Apex, which outperforms frontier models in customer service tasks, and is the foundation of the AI Group's full stack approach to AI.

We’re Particularly Interested In Engineers Who …

https://www.linkedin.com/jobs/view/4403460327

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: AI Infrastructure Engineer – Mark Allen Evans

Hi,

Fin's push into deeper reasoning and multi-provider orchestration at scale is exactly the kind of infrastructure problem I've been living in.

At AiAssist Secure, I architected a BYOK multi-provider AI platform spanning 11+ model providers — with routing logic, governance controls, and enterprise-facing control surfaces built from the ground up. That's the same layer where Fin's reliability gets won or lost.

I'm exploring what's next and Intercom's infrastructure team stood out immediately.

If you're not the right person to pass this to, no worries at all — happy to connect with whoever owns this space, even just for 20 minutes.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4403460327 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI & Automation Engineer, Endpoint Systems- Remote

70% fit

BeyondTrust · Canada (Remote) · Remote ' · linkedin_jobs

ROLE

BeyondTrust is a place where you can bring your purpose to life through the work that you do, creating a safer world through our cybersecurity SaaS portfolio.

Our culture of flexibility, trust, and continual learning means you will be recognized for your growth, and for the impact you make on our success. You will be surrounded by people who challenge, support, and inspire you to be the best version of yourself.

The Role

We’re looking for an Endpoint Systems Engineer with a passion for AI & automation, within modern device management. In this role, you’ll design, deploy, and continuously improve endpoint solutions that are intelligent, scalable, and secure. You’ll move beyond reactive support by leveraging automation, data, and AI-driven tools to enhance user experience and operational efficiency across the enterprise. This is a hands-on engineering role for someone who enjoys building smart systems, eliminating repetitive work, and shaping the future of endpoint management.

What You’ll Do

Design, deploy, and manage modern endpoint environments across Windows & macOS. Implement and maintain endpoint management platforms (e.g., Intune, Jamf) with a strong automation-first mindset Develop automation workflows using scripting, APIs, and orchestration tools to reduce manual effort and improve reliability Integrate AI-powered tools for endpoint analytics, remediation, security, and user experience optimization Build self-healing and self-service solutions for common endpoint …

https://www.linkedin.com/jobs/view/4380578483

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: AI & Automation Engineer role — passing this along if it lands wrong

Hi,

BeyondTrust's push to layer AI into PAM and endpoint privilege enforcement is one of the more technically interesting bets in identity security right now — the surface area is real and the stakes are high.

I'm a fit worth a look. At AiAssist Secure, I designed secure runtime separation between customer-facing APIs and isolated execution services — signed internal requests, allowlisted access, auditable actions. That's the same control-plane thinking your endpoint AI work demands.

Four years in AI systems engineering, deep in Python, Node.js, and platform architecture.

Portfolio: AiAssist.net | GitHub: github.com/aiassistsecure

Happy to connect — or if hiring runs through someone else, feel free to forward.

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4380578483 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI Engineer (Remote)

Lumenalta · Canada (Remote) · Remote ' · linkedin_jobs

70% fit

ROLE

At Lumenalta, we partner with forward-thinking organizations to build technology solutions that scale, delight users, and accelerate business growth. Our global teams bring curiosity, commitment, and technical excellence to every project. We value transparency, autonomy, and impact—empowering every team member to do their best work.
We're seeking experienced AI Engineers to join a high-growth fintech prospect focused on back-office workflow automation. This role is an opportunity to design and deploy production-grade AI systems for community banks, mortgage, and insurance institutions while shaping a flexible, modern architecture alongside an accomplished executive team.
Future Opportunity Role – Talent PipelineThis is a future opportunity role. We continuously meet talented engineers to support upcoming client projects. While there may not be an immediate opening, qualified candidates may be considered for future engagements.
What You'll Be DoingDesign, build, and deploy AI models into production (not just prototypes or wrappers).Develop and maintain backend Python services that power intelligent automation workflows.Implement enterprise-scale data pipelines for ingesting and processing financial and regulatory documents, with automated data quality checks and monitoring.Train and optimize models on banking rules, compliance standards, and regulatory datasets.Build intelligent routing systems (e.g., auto-approval for simple loans/accounts, escalation to humans for complex ca…

https://www.linkedin.com/jobs/view/4401018556

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: AI Engineer role at Lumenalta — passing this along if it lands wrong

Hi,

Lumenalta's push into back-office workflow automation for fintech caught my attention — exactly the kind of applied AI work I want to be doing.

One thing that might be relevant: I architected a multi-provider AI platform at AiAssist Secure with routing, governance, and enterprise control surfaces across 11+ model providers — the kind of production-grade system that maps closely to what this role is describing.

Portfolio: AiAssist.net | GitHub: github.com/aiassistsecure

If you're not the right person to loop in, happy for you to forward this — otherwise I'd love 20 mins with whoever owns the AI hiring.

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4401018556 · captured: 2026-04-18 · fit_reason: keyword + remote match

Infrastructure Automation AI Developer

Nokia · Kanata, ON (On-site) · On-site · linkedin_jobs

70% fit

ROLE

Job Description

In late 2024, we introduced a new infrastructure automation application called EDA.EDA is an award-winning cloud-native, scalable, extensible platform for automating the management of networks and data centers. Join a growing team of talented engineers who are building an innovative new platform for the next generation of infrastructure that drives the modern world.

We believe in Subject Matter Expertise and code ownership. We keep the process simple so that developers can focus on coding. With minimum overhead, developers interact directly with Product Managers, peer designers, and QA. If you are proud of your code, you belong on our team.

DCN2025

How You Will Contribute And What You Will Learn

Work with product line managers to gather requirements and work with team members to come up with efficient software design to fit the requirements. Develop, unit-test, and integrate software features for EDA's AI features.Work with QA engineers to trouble-shoot complex problems to find best possible resolution post root cause analysis. Willing to take end-to-end responsibility of a file/module (deal with new requirements, prioritize work, maintain quality, support and maintain).

Key Skills And Experience

Minimum Bachelor's degree in Computer/Software Engineering, Computer Science, equivalent, or better. Minimum 8-10 years of hands-on experience in the industry. Good working knowledge of golang, especially for implementing gRPC based servers.Familiarity with la…

https://www.linkedin.com/jobs/view/4341145287

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: EDA team @ Nokia — infrastructure automation background, worth a look?

Hi,

Nokia's EDA launch caught my attention — cloud-native, event-driven network automation at that scale is exactly the problem space I've been working in.

I'm Mark. Most recently I architected a multi-provider AI platform with runtime isolation, signed internal request flows, and auditable control surfaces across 11+ model providers — the kind of governance and execution separation that maps directly to what EDA is building.

4 years in platform and AI systems engineering, strong in Python, Node.js, and TypeScript.

If you're not the right person to loop in, no worries — happy if you can point me that direction. Otherwise, I'd love 20 mins with whoever owns this.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4341145287 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer

Monzo Bank · London, England, United Kingdom (Remote) · Remote ' · linkedin_jobs

70% fit

ROLE

Ø=Þ€ We’re on a mission to make money work for everyone.

We’re waving goodbye to the complicated and confusing ways of traditional banking.

After starting as a prepaid card, our product offering has grown a lot in the last 10 years in the UK. As well as personal and business bank accounts, we offer joint accounts, accounts for 16-17 year olds, a free kids account and credit cards in the UK, with more exciting things to come beyond. Our UK customers can also save, invest and combine their pensions with us.

With our hot coral cards and get-paid-early feature, combined with financial education on social media and our award winning customer service, we have a long history of creating magical moments for our customers!

We’re not about selling products - we want to solve problems and change lives through Monzo 'dþ

Ø=ÜÍLondon / UK Remote | Ø=Ü°Senior Engineer £95,000 - £130,000 + Incentive Awards tied to your performance Benefits | Technology - Engineering |

About Platform Engineering At Monzo

The Platform Collective builds and maintains the infrastructure, tools and processes that sets the rest of Monzo technology teams up for success. We work on a wide range of shared infrastructure, services and engineer tooling.

The Infrastructure Platform team sit within the Platform Collective and are responsible for running the Kubernetes clusters, service mesh and networking that all of Monzo's microservices run on. Our team is made up of backend engineers and cloud infrastructure experts, …

https://www.linkedin.com/jobs/view/4386839381

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Platform Engineer role @ Monzo

Hi,

Monzo's push into pensions, investing, and multi-product infrastructure is a genuinely interesting engineering problem — the platform complexity behind that kind of expansion is real.

I'm a platform engineer with 4 years building in that space. Most recently, I designed runtime separation between customer-facing APIs and isolated execution services using signed internal requests, allowlisted access, and auditable action logs — the kind of control surface that scales across products without becoming a liability.

I'd love 20 minutes with whoever owns platform hiring. If that's not you, happy for you to pass this along.

GitHub: github.com/aiassistsecure
Portfolio: AiAssist.net

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4386839381 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer | Remote

Crossing Hurdles · United Kingdom (Remote) · Remote ' · linkedin_jobs

70% fit

ROLE

Position: DevOps / Platform EngineerType: Hourly contractCompensation: $70–$150 per hourLocation: RemoteCommitment: Ongoing, flexible part-time or project-based engagement
Role ResponsibilitiesBuild and manage CI/CD pipelines and deployment workflowsDevelop and maintain cloud infrastructure across AWS, GCP, or AzureImplement containerization and orchestration using Docker and KubernetesWork independently to support engineering and platform reliabilityCollaborate with teams and communicate technical decisions clearly
RequirementsStrong experience in DevOps or platform engineeringExperience with CI/CD pipelines and automation workflowsFamiliarity with cloud infrastructure (AWS, GCP, or Azure)Experience with containerization and orchestration tools (Docker, Kubernetes)Strong communication skills and ability to work independently in a remote environment
Application Process (Takes 20 Mins)Upload your resume.Complete an interview.Submit a short form.

https://www.linkedin.com/jobs/view/4385252554

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Re: Platform Engineer | Remote at Crossing Hurdles

```json
{
  "subject": "Platform Engineer role — Mark Allen Evans",
  "body": "Hi,

Noticed Crossing Hurdles is placing contract platform engineers for clients who need on-demand infrastructure depth rather than full-time hires — that model makes a lot of sense right now.

I'm Mark. At AiAssist Secure, I designed secure runtime separation between customer-facing APIs and isolated execution services — signed internal requests, allowlisted access, auditable actions. The kind of reliability and security work that maps directly to what your clients likely need from a platform hire.

I work well independently, communicate clearly across technical and non-technical teams, and I'm open to flexible contract structures.

If you're not the right person to send this to, no worries — happy if you can point me in the right direction. Otherwise, would love 20 mins.

Mark Allen Evans
https://AiAssist.net"
}
```

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4385252554 · captured: 2026-04-18 · fit_reason: keyword + remote match

Staff Platform Engineer, Canada

Ashby · Montreal, QC (Remote) · Remote ' · linkedin_jobs

70% fit

ROLE

We’re looking for a curious, rigorous, problem-hungry platform engineer (who codes!) to carry the ball as we bring Ashby to the big leagues. Ashby builds software that lets talent teams build an efficient, delightful, respectful hiring process. Similarly, you’re an engineer who wants to build a “paved road” that excellent engineering teams can safely take to the moon and back.

We have notable customers like Notion, Linear, Shopify, and Snowflake. Our growth and retention metrics are best-in-class among our peers: we have tens of millions in ARR, growing >100% year over year, over 2500 customers, very low churn, and many years of runway. We’ll share more details once we meet, but you now probably have a good idea as to why we're hiring for this role Ø=Þ .

We’ve listed this role twice: as a Platform Engineer and Site Reliability Engineer – our team does both, and we are open to candidates who lean towards one or the other.

About The Role And How We Work

Hi Ø=ÜK I’m Colin, Head of EMEA Engineering. I’ve spent a number of years leading engineering teams in startups, and that has always included being close to infrastructure teams - no matter what name they’ve worn (SRE, infrastructure, platform, etc). I’ve got my hands dirty building the initial infrastructure for startups and know the value a talented infrastructure engineer brings. The rigour, the discipline, the peace and quiet when everything just hums along.

Our infrastructure is in a good place for now. Nothing is static. …

https://www.linkedin.com/jobs/view/4323336839

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Staff Platform Engineer role — passing this along if it lands wrong

Hi,

Ashby's push into enterprise-grade platform infrastructure — paved roads for fast-scaling engineering teams — caught my attention. That's exactly the kind of problem I've been living in.

At AiAssist Secure, I designed secure runtime separation between customer-facing APIs and isolated execution services using signed internal requests, allowlisted access, and auditable action trails — the kind of trust boundary work that matters when you're serving customers like Notion and Snowflake.

I'm based in Canada and genuinely curious about what Ashby's platform team is building next.

If you're not close to hiring for this, happy for you to point me toward whoever is — or just ignore this entirely.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4323336839 · captured: 2026-04-18 · fit_reason: keyword + remote match

Founding DevRel Engineer

65% fit

Nozomio · San Francisco, CA, US · On-site · ycombinator

Ø=Ü° $120K - $170K

ROLE

Founding DevRel Engineer at Nozomio | Y Combinator Nozomio index anything. search it with ai. Founding DevRel Engineer $120K - $170K • 0.05% - 0.25% • San Francisco, CA, US Job type Full-time Role Operations Experience Any (new grads ok) Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Arlan Rakhmetzhanov Founder Arlan Rakhmetzhanov Founder About the role TL;DR We are hiring a founding DevRel Engineer for Nozomio, the context augmentation layer for AI agents. You will write content, documentation, and you will be the face of Nozomio for developers. You will define how AI engineers think about context, MCP servers, and agentic coding. You will own DevRel end to end and over time build the function around you. About Nia and Nozomio Nozomio is applied product and research lab. Our flagship product Nia is a context augmentation layer for AI agents, primarily coding agents. It indexes entire codebases, documentation, and external sources so agents always have the right context instead of guessing. Nia plugs in as an MCP server into tools like Cursor, Cline, and other coding agents, giving them a persistent, searchable, up-to-date knowledge base. Today, agents are very good at generating code but still bad at knowing which repo, package, or doc to look at. Context, not generation, is the bottleneck. Nia exists to solve that bottleneck so agents reach near superhuman code intelligence. Nozomio is backed by Y Combinator and rece…

https://www.ycombinator.com/companies/nozomio/jobs/SNcQ2l0-founding-devrel-engineer

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Founding DevRel – Mark Allen Evans

Hey,

Nia's approach to context augmentation as an MCP layer is the right bet — context quality really is the bottleneck, and most teams are still duct-taping retrieval together.

I've been building in this space: recently shipped agent and app-generation workflows with tool routing, memory-aware patterns, and developer-facing API/SDK surfaces — pretty close to what a founding DevRel role here would touch daily.

Portfolio: AiAssist.net | GitHub: github.com/aiassistsecure

If you're not the right person to ping on this, happy to be pointed in the right direction — otherwise, would love 20 mins with whoever owns the search.

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:86596 · captured: 2026-04-18 · fit_reason: keyword + remote match

Staff AI Engineer - USA

65% fit

Podium · Lehi, UT, US / Remote (US) · Remote ' · ycombinator

Ø=Ü° $190K - $215K

ROLE

Staff AI Engineer - USA at Podium | Y Combinator Podium Get more leads. Make more money. Staff AI Engineer - USA $190K - $215K • Lehi, UT, US / Remote (US) Job type Full-time Role Engineering, Full stack Experience 11+ years Visa US citizen/visa only Apply to Podium and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role At Podium, our mission is to arm every local business with a complete platform and outcome-driven AI employees that convert leads into real, paying customers. Every day, millions of workers use our AI lead conversion and communication platform to help them get more leads and make more money. Our work and focus on helping local businesses thrive has been recognized across the industry, including Forbes' Next Billion Dollar Startups, Forbes' Cloud 100, the Inc. 5000, and Fast Company's World's Most Innovative Companies. At Podium, we believe in fostering a culture that thrives on hiring and developing exceptional talent. Our operating principles serve as a compass, guiding daily behavior and decision-making, and ensure we hire people who will thrive at Podium. If you resonate with our operating principles and are energized by our mission, Podium will be a great place for you! Staff AI Engineer The Role We are looking for a talented Staff AI Engineer to help build and scale our AI Agent platform—a powerful system that autonomously interacts with customers, handling millions of conversations every month. In this role, …

https://www.ycombinator.com/companies/podium/jobs/qgKxAEM-staff-ai-engineer-usa

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: Staff AI Engineer – worth a quick look

Hey —

Podium's push toward outcome-driven AI employees for local businesses — not just chat widgets but actual revenue infrastructure — is one of the more grounded bets I've seen in this space.

I built something in the same vein: a multi-provider AI platform with routing, governance, and enterprise control surfaces across 11+ model providers — essentially the plumbing that makes agent behavior predictable at scale.

That's the core of what your Staff AI Engineer role seems to need, and I'd love 20 minutes with whoever owns that search.

If that's not you, no pressure — happy if you can point me the right direction.

Mark Allen Evans
https://AiAssist.net

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:88046 · captured: 2026-04-18 · fit_reason: keyword + remote match

Ingénieur IA Ops

Devoteam · Île-de-France, France (Hybrid) · On-site · linkedin_jobs

65% fit

ROLE

Devoteam est un leader du conseil en stratégie digitale, plateformes technologiques, cybersécurité et business transformation. Avec 10 000 collaborateurs en Europe, Moyen-Orient et Afrique, Devoteam accompagne les entreprises en mettant depuis plus de 30 ans la Technologie au service de l'Homme.Les systèmes d'information sont devenus des actifs critiques. Pourtant, trop souvent, leur gestion reste invisible, sous-valorisée, ou fragmentée. Sans une base solide, aucune transformation digitale ne tient dans la durée. Sans un run fiable, il n'y a pas de performance durable.Notre mission : garantir la paix d'esprit technologique, en assurant la continuité, la fiabilité et la performance des environnements IT. L'équipe de TopTeam est ancrée dans l'opérationnel, portée par la rigueur, la proximité et la culture du service.Nos experts -Service Delivery Manager, Chef de Projet Production, Ingénieur de Production Devops, AIOps, Ingénieurs & AdministrateursSystèmes et Réseaux, Ingénieur Digital Workplace, Support technique N2/N3 et VIP interviennent sur site ou à distance, avec un delivery piloté par les SLA et l'amélioration continue. Nous industrialisons les opérations, fiabilisons les environnements et sécurisons le quotidien.Notre engagement : tenir le run pour libérer la transformationPourquoi nous rejoindre ? Nous sommes une IA-Driven Company : Né lors des prémices de la révolution tech, Devoteam est le partenaire de confiance qui a accompagné ses clients vers l'évolution d'Intern…

https://www.linkedin.com/jobs/view/4403416074

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

Subject: AI Ops role at Devoteam — passing this along if it lands wrong

Hi,

Devoteam's push to make AI infrastructure visible and governable inside enterprise SI stacks caught my attention — it's exactly the gap most orgs are still fumbling through.

I'm Mark. At AiAssist Secure, I architected a multi-provider AI platform with routing, governance, and enterprise-grade control surfaces across 11+ model providers — the kind of ops layer your team is likely building around.

Four years in AI systems engineering, comfortable across Python, Node.js, and API-layer design.

If you're close to the team hiring for this, I'd love 20 minutes with whoever owns it. Happy to forward my portfolio: AiAssist.net

Mark Allen Evans

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4403416074 · captured: 2026-04-18 · fit_reason: keyword + remote match

Cloud Engineer - AI Trainer

DataAnnotation · United Kingdom (Remote) · Remote ' · linkedin_jobs

65% fit

ROLE

Join the DataAnnotation team and contribute to developing cutting-edge AI systems, while enjoying the flexibility of remote work and setting your own schedule.

We are looking for proficient programmers to help advance AI development. As a member of DataAnnotation's coding team, you'll be part of a growing community of over 100,000 professionals — including front-end, back-end, full-stack, machine learning, and other engineers — who are driving real-world impact in AI development.

Our platform offers an engaging blend of flexibility and challenge: you'll work closely with state-of the art AI models to take on programming tasks that include creating and solving challenging coding problems, building beautiful apps with rich functionality, and synthesizing insights through data analysis and visualization. Your work directly contributes to refining intelligent systems that learn, adapt, and evolve. Some team members fit this work alongside a full-time role, while others treat it as their primary focus, choosing projects and schedules that align with their availability and goals.

To get started, once you sign up for an account, you'll take a short assessment (this serves as our version of an interview). If you pass that assessment, you'll receive an email confirmation, and paid work will become available to you through our platform.

Benefits

Fully remote: work from anywhere in the US, Canada, UK, Ireland, Australia, and New Zealand.Flexible schedule: choose which projects you …

https://www.linkedin.com/jobs/view/4397395531

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4397395531 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer

incident.io · London, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

65% fit

ROLE

About Incident.io

incident.io is the leading AI incident response platform, built to help teams dramatically reduce incident response time and improve reliability. We bring together on-call, incident response, AI SRE, and status pages in a single platform, giving teams everything they need to respond quickly, reduce downtime, and keep customers in the loop.

Since launching in 2021, we've helped over 1,500 companies, including Netflix, Airbnb, and Block, run more than 500,000 incidents. Every month, tens of thousands of responders across Engineering, Product, and Support use incident.io to restore services faster, stay aligned under pressure, and focus on building what matters.

We're a fast-growing, highly ambitious team that cares deeply about our customers, product quality, and making it magic. We've raised $100M from Index Ventures, Insight Partners, and Point Nine, alongside founders and executives from world-class technology companies.

The Team

Our Platform Engineers are the people behind the curtain making sure everything runs smoothly. They own the infrastructure, deployment systems, and CI/CD pipelines that keep our product engineers moving fast without breaking things. If product engineers are the chefs, the platform team is the kitchen; designing, maintaining, and scaling the tools that make everything possible.

The current team consists of two highly capable engineers. As we grow, this role will play a key part in building and maintaining infrastructure to sup…

https://www.linkedin.com/jobs/view/4350971911

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4350971911 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer - Cloud

65% fit

EDF (UK) · Gloucester, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

About The Role

Ready to make an impact by shaping secure, modern cloud solutions that support the systems keeping our nuclear fleet running safely and efficiently? Want to grow your skills in Azure engineering while contributing to vital national infrastructure? At EDF, Success is Personal – and with us, your growth journey is yours to shape.

Working in our Nuclear IT Services team, you'll help deliver the capabilities supporting the transition towards An Electric Britain. You'll build trusted cloud services that support critical operations and make a meaningful contribution every day.

This role can be based anywhere across our fleet, working to a hybrid model, giving you the flexibility to balance time at home with time on site when needed.

What You'll Be Doing

Deploying, managing and enhancing Azure servicesBuilding, deploying, and maintaining CI/CD pipelines using Azure DevOps or GitHub Actions.Administration of Windows and Linux systems in AzureSupporting automation through PowerShell, Azure CLI and related toolsMonitoring performance, cost, governance and security across cloud environments

Who You Are

We're looking for a Cloud Platform Engineer who thrives on solving problems, working collaboratively, and delivering effective technical outcomes. Do you offer…?

Demonstrated capability with Infrastructure as Code (IaC)Proficiency in Azure ServicesExpertise in Windows and Linux system administrationHighly proficient with cloud automation tools such as Azure DevOps, …

https://www.linkedin.com/jobs/view/4400454475

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400454475 · captured: 2026-04-18 · fit_reason: keyword + remote match

LLM Training Engineer (RLHF & Fine-tuning)

65% fit

Confidential · United Kingdom (Remote) · Remote ' · linkedin_jobs

ROLE

About Us
Leading global fintech platform serving millions of users, top-ranked by transaction volume.We're pioneering AI-native systems where LLMs, intelligent agents, and automation drive core capabilities—not just features.

The Role
Seeking engineers to build vertical AI agents for financial technology domains.Train, fine-tune, and deploy LLMs (using SFT/RLHF) as production-ready systems, including:
Fintech AI agents & copilotsDomain Q&A and knowledge toolsSpecialized financial insight engines
Collaborate with product/engineering teams to ship AI at scale.

What You'll Work On
LLM training/fine-tuning (SFT, RLHF, distillation) for financial use casesRAG systems & knowledge integrationModel optimization for real-time productionAI agent development & efficiency pipelines

Who You Are
Hands-on LLM/deep learning experience (PyTorch/Transformers)Built domain-specific AI (fintech/finance preferred)Strong Python & engineering fundamentalsMandarin & English fluency required
Nice to Have
Fintech/trading/quant backgroundScaled AI deploymentsAgentic AI systems
Why Join
Shape AI for a global top platformReal-world impact at massive scaleFully remote (GMT+8 hours focus)

https://www.linkedin.com/jobs/view/4391528886

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4391528886 · captured: 2026-04-18 · fit_reason: keyword + remote match

Infrastructure/SRE Engineer - Qualcomm - Remote, UK

65% fit

Qualcomm · London, England, United Kingdom (Remote) · Remote ' · linkedin_jobs

ROLE

Company

Qualcomm Technologies International Ltd

Job Area

Engineering Group, Engineering Group > Software Engineering

General Summary

We are looking for an Infrastructure/SRE Engineer to help build and evolve the cloud platform that powers Edge Impulse. You will work alongside the infrastructure team to manage AWS infrastructure with Terraform, maintain Kubernetes clusters with Karpenter and Cilium, build CI/CD pipelines in Buildkite, and drive GitOps deployments through ArgoCD. Our platform runs a mix of CPU and GPU workloads to support ML model training and deployment, making infrastructure decisions impactful on both cost and performance. This role is part of the foundation that enables all engineering teams to ship reliably and securely at scale.

Preferred Qualifications

4+ years in SRE, DevOps, or Platform Engineering rolesStrong AWS experience (EKS, RDS, S3, IAM, VPC)Proficiency with Terraform in production environmentsSolid Kubernetes knowledge including networking, RBAC, and autoscaling (Karpenter)Experience with Cilium or eBPF-based networking, including network policies, observability with Hubble, and service mesh conceptsExposure to GPU-accelerated workloads on Kubernetes is a plusExperience with ArgoCD or similar GitOps workflowsCI/CD pipeline experience (Buildkite or similar)Proficiency in TypeScript, Python, or BashFamiliarity with Datadog, HashiCorp Vault, PostgreSQL, and RedisStrong communicator comfortable on a fully distributed team

Roles/Responsibil…

https://www.linkedin.com/jobs/view/4389593535

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4389593535 · captured: 2026-04-18 · fit_reason: keyword + remote match

Staff AI Engineer (Global Security)

RBC · Toronto, ON (Remote) · Remote ' · linkedin_jobs

65% fit

ROLE

Job Description

What is the opportunity?

The Staff AI Engineer (Global Security) will be part of Global IT Risk's critical modernization initiative, designing and building advanced AI models that power our IT Risk & Control Assessment platform. Your models will provide data-driven Inherent and Residual Risk Ratings across all RBC applications globally, scaling our assessment capability from today's manual, subset approach to full automation across all domains. These Assessments aim to provide a data-driven Inherent and Residual Risk Rating for every application in RBC across domains of Risk. To conduct a Risk Assessment and their supporting calculations, several key data points are required, including Control Assessments, IT Issues Management findings, Control Environment details and Control Testing results. These data points will drive effectiveness and reliability in the calculation of Residual Risk. The Residual Risk and associated control gaps will then be aggregated to formulate action plans and a risk score for both domain owners and Applications to executives. This solution will scale out our assessment capability across all appcodes globally and domains and will direct leadership decisions on how to reduce the impact and likelihood of regulatory, financial or reputational risks to RBC.

What will you do?

Design and build AI models that automate control testing, risk assessment, and anomaly detection across IT domains, serving as consultant on the estate of models a…

https://www.linkedin.com/jobs/view/4399212375

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399212375 · captured: 2026-04-18 · fit_reason: keyword + remote match

Développeur(se) logiciel – AI-first

65% fit

Kampwise · Quebec, Canada (Remote) · Remote ' · linkedin_jobs

ROLE

Développeur(se) logiciel – AI-first(Autodidacte, orienté efficacité, passionné par les outils AI)
À propos de Kampwise:Kampwise est une startup technologique québécoise en pleine croissance qui modernise la gestion et la réservation de campings.On vient de compléter une première ronde de financement de 1M$, et on entre dans une phase clé : Structurer notre produit, accélérer notre développement et bâtir notre équipe technologique à l'interne.Nous sommes basés dans la région de Québec et nous recherchons une personne qui habite dans la région ou à proximité.Aujourd'hui, notre développement est assuré en collaboration avec un partenaire externe.Demain, on veut construire une équipe forte à l'interne.Et cette première embauche va jouer un rôle central dans cette transition.
Ton rôle chez nous:On ne cherche pas juste un dev.On cherche quelqu'un qui veut bâtir, optimiser, et faire évoluer notre façon de développer.Tu seras au cœur de notre produit et de notre évolution technologique :Tu travailles en collaboration avec notre partenaire technologique (équipe externe)Tu comprends, analyses et challenges le code existantTu participes activement aux décisions techniques et aux améliorationsTu prends en charge les tests, la qualité et la stabilité du produitTu identifies et corriges les enjeux avant qu'ils arrivent en productionTu optimises les processus de développementEt surtout :Tu mets en place et exploites les outils AI pour rendre le développement plus rapide, plus intelligent et…

https://www.linkedin.com/jobs/view/4399232752

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399232752 · captured: 2026-04-18 · fit_reason: keyword + remote match

Staff AI Engineer - Grafana Ops, AI/ML | Canada | Remote

65% fit

Grafana Labs · Canada (Remote) · Remote ' · linkedin_jobs

ROLE

Grafana Labs is a remote-first, open-source powerhouse. There are more than 20M users of Grafana, the open source visualization tool, around the globe, monitoring everything from beehives to climate change in the Alps. The instantly recognizable dashboards have been spotted everywhere from a NASA launch and Minecraft HQ to Wimbledon and the Tour de France. Grafana Labs also helps more than 3,000 companies -- including Bloomberg, JPMorgan Chase, and eBay -- manage their observability strategies with the Grafana LGTM Stack, which can be run fully managed with Grafana Cloud or self-managed with the Grafana Enterprise Stack, both featuring scalable metrics (Grafana Mimir), logs (Grafana Loki), and traces (Grafana Tempo).

We’re scaling fast and staying true to what makes us different: an open-source legacy, a global collaborative culture, and a passion for meaningful work. Our team thrives in an innovation-driven environment where transparency, autonomy, and trust fuel everything we do.

You may not meet every requirement, and that’s okay. If this role excites you, we’d love you to raise your hand for what could be a truly career-defining opportunity.

This is a remote opportunity and we would be interested in applicants from Canada time zones only at this time.

Staff AI Engineer

The Opportunity

At Grafana, we build observability tools that help users understand, respond to, and improve their systems – regardless of scale, complexity, or tech stack. The Grafana AI teams play a…

https://www.linkedin.com/jobs/view/4346487893

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4346487893 · captured: 2026-04-18 · fit_reason: keyword + remote match

Lead Platform Engineer - Global AI Platform

65% fit

Manulife · Toronto, ON (Hybrid) · On-site · linkedin_jobs

ROLE

As a Lead Platform Engineer, you will help design, build and deliver the foundational services that enable secure, scalable, and reusable AI and multi-agentic workflow solutions across the organization. You will work across cloud, distributed systems, orchestration frameworks, vector/memory stores, agent workflows, and safety/guardrail layers to ensure the platform is secure, scalable, and enterprise ready.

Position Responsibilities

Builds and maintains high-performance, fault-tolerant, secure, and scalable AI platform services and abstractions that support diverse AI solutions with automation-first delivery.Designs, builds, and maintains the technology platform's features and infrastructure, including hardware, software, and network components.Build scalable microservices and event-driven pipelines for model training/inference using Akka Streams and Cluster Sharding.Integrate AdaptiveML workflows for continuous/online learning, feature stores, model registries, and A/B experimentation.Implement AI Foundry components for orchestration, feature engineering, model deployment, and governance.Develop reusable reference patterns, inner-source components that meet reliability, security, and compliance standards.Implement shared runtimes for multi agent coordination, state management, memory persistence, and messaging.Design interoperable APIs/SDKs used by data scientists and developers to build agent powered applications.Maintain and improve CI/CD pipelines and developer toolchai…

https://www.linkedin.com/jobs/view/4391879351

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4391879351 · captured: 2026-04-18 · fit_reason: keyword + remote match

Développeur déployé I, IA générative, Google Cloud

65% fit

Google · Montreal, QC (On-site) · On-site · linkedin_jobs

ROLE

Cette offre d'emploi est pour un nouveau poste vacant.

Remarque : En postulant à cette offre, vous pourrez indiquer où vous préférez travailler parmi les propositions suivantes : Montreal, QC, Canada; Toronto, ON, Canada.Qualifications minimales:

Baccalauréat en science, en technologie, en ingénierie ou en mathématiques, ou expérience pratique équivalenteExpérience de la création et de la livraison de solutions de qualité de production axées sur l'IA à des clients externes ou internes à l'aide de Python, de Typescript ou de langages similairesExpérience en IA appliquée, avec un accent sur la création de systèmes autour de modèles préentraînés (p. ex. l'ingénierie des requêtes, l'affinement, la GAR, l'orchestration des interactions entre les modèles et les outils externes pour fournir des solutions)Expérience en matière de conception, de déploiement ou de gestion de solutions sur une plateforme infonuagique (p. ex. Google Cloud Platform)

Qualifications souhaitées:

Maîtrise ou doctorat en IA, en informatique ou dans un domaine technique connexeExpérience de la mise en œuvre de systèmes multiagents à l'aide de cadres (p. ex. LangGraph, CrewAI ou ADK de Google) et de modèles tels que React, l'autoréflexion et la délégation hiérarchiqueConnaissance des mesures « natives des LLM » (p. ex. jetons/s, coût par requête) et des techniques d'optimisation de la gestion des états et du traçage granulaire

null

En tant que développeur déployé d'IA chez Google Cloud, vous serez un dével…

https://www.linkedin.com/jobs/view/4402723677

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4402723677 · captured: 2026-04-18 · fit_reason: keyword + remote match

Staff Applied AI Engineer — Auggie Platform - CANADA (Remote)

65% fit

Luxury Presence · Greater Montreal Metropolitan Area (Remote) · Remote ' · linkedin_jobs

ROLE

Luxury Presence is building the AI growth platform for real estate. Backed by Bessemer Venture Partners and other top investors, we're a Series C company on track to hit $100M in annual recurring revenue in the next six months. More than 87,000 real estate professionals, including over 30% of the WSJ Real Trends top 100 agents in the United States, use us to run and grow their business.

The Opportunity

We don't just use AI. We build with it, ship with it, and think with it. We're in the top 1% of companies applying AI effectively, not just to our products, but to how we build them. With an unlimited budget for Anthropic tokens, our engineers use AI agents to write, review, and ship production code every day. We're building toward a world where humans design systems and AI builds them, and we're already further along that path than almost anyone else.

As a Staff Software Engineer, you'll be a technical leader at the center of this transformation. You'll shape platform architecture, drive AI-powered product delivery, and raise the bar for what a small, AI-augmented engineering team can accomplish. This is a role for someone who already lives in their terminal with AI, who has strong opinions about how AI changes system design, and who wants to help define what a next-generation engineering organization looks like.

What You'll Do as a Staff Engineer on Our Auggie Platform Team

Here are examples of the kinds of initiatives you'd work on today:

Multi-agent orchestration and …

https://www.linkedin.com/jobs/view/4392488619

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4392488619 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior AI/ MLOps Engineer H/F

60% fit

Devoteam Data Driven · Nantes, Pays de la Loire, France (Hybrid) · On-site · linkedin_jobs

ROLE

Ø>Ý¸ &Bþ Présentation de l’équipe

Devoteam Nantes, avec ses 160 collaborateurs, est un acteur clé de la transformation digitale dans la région nantaise. Au sein de l’agence, l’entité Data Driven a pour ambition d’aider ses clients à transformer la valeur de leurs données en succès durable et d’accélérer leur adoption de l’IA au sein de leurs équipes métiers.

Nous tirons parti à la fois de notre expertise sur les plateformes technologiques les plus innovantes et de notre connaissance métier avec la meilleure combinaison d’experts pour accompagner nos clients dans leur quête d’agilité, de compétitivité et de croissance durable.

Pour contribuer à ces évolutions, nous sommes à la recherche d’un(e) consultant(e) Senior AI Engineer souhaitant rejoindre un collectif d’une dizaine de consultant(e)s spécialisé(e)s en data engineering à Nantes.

Description Du Poste

Ø=Üª Vos missions :

Supporter le processus de développement machine learning de bout en bout pour concevoir, créer et gérer des logiciels reproductibles, testables et évolutifs. Travailler sur la mise en place et l’utilisation de plateformes IA en s’appuyant sur les concepts de ML/LLMOpsApporter à nos clients des best practices en termes d’organisation, de développement, d’automatisation, de monitoring, de sécurité. Expliquer et appliquer les best practices pour l’automatisation, le testing, le versioning, la reproductibilité et le monitoring de la solution IA déployée. Encadrer et superviser les consultant(es) juniors i.…

https://www.linkedin.com/jobs/view/4390149813

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4390149813 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior AI Engineer (x/f/m)

Doctolib · Nantes, Pays de la Loire, France (On-site) · On-site · linkedin_jobs

60% fit

ROLE

Your ImpactWe are looking for a Senior AI Engineer to join the Customer Care Contact team in the Business Performance domain.Your mission will be to build AI-powered tools that revolutionize how healthcare professionals deliver care. You will work on Doctolib Assistant, our AI personal assistant designed to assist doctors and medical staff in their daily work—from answering complex medical questions to automating administrative tasks with accuracy, reliability, and safety.Working in the tech team at Doctolib means building innovative products and features to improve the daily lives of care teams and patients.What you'll build

Your responsibilities include but are not limited to: Design, ship, and operate production-grade agentic systems powering Doctolib Assistant, our AI personal assistant that addresses key challenges for healthcare professionalsBuild and improve AI agents and our orchestration systemsPartner with data scientists to design and implement evaluation pipelines, security/safety guardrails, and metrics—while taking ownership of operational AI projects as data science teams focus on exploratory researchDevelop and maintain core agentic capabilities including memory, reasoning orchestration, tool adapters, and multi-agent coordination, ensuring strong observability, fallback mechanisms, and fail-safe behaviorsOwn performance, observability and security: monitoring dashboards, SLO/SLI, and security postureEnsure privacy and security by design, with a particular fo…

https://www.linkedin.com/jobs/view/4378828506

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4378828506 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior AI/LLM Engineer

Criteo · Paris, Île-de-France, France (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

What You'll Do:

We are looking for an experienced AI Engineer to join the Creator ML team.

We apply state of the art Machine Learning to optimize all possible aspects of the construction and serving of our banners. In this role, you will work across the full lifecycle of large language model applications: from dataset creation and evaluation to agent APIs, observability, and production-grade LLM features.

More broadly, within the Creator group, we are on a mission to revolutionize the creative capabilities offered to our clients across our commerce media platform. You will play a crucial role in shaping the landscape of Internet advertising, and you will have the opportunity to work on projects that impact billions of users worldwide.

Your main missions

Develop LLM-powered features which include optimizing prompts, model selection, and inference strategies for quality, latency, and cost. You will be implementing end-to-end observability capabilities and agent executions. As well, you will have the opportunity to shape the development of infrastructure for Agents and MCP-based agent interfaces to support Agentic Creative workflows. Ensure systems are robust, testable, and reproducible. You will be contributing to internal tooling for experimentation, evaluation, and monitoring of Large Language Models. Challenge the status quo of our creative tooling for creation and optimization by leveraging more recent advancements in Generative AI. You will as well partner closely wit…

https://www.linkedin.com/jobs/view/4390813729

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4390813729 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI Engineer

DEPT® · Berling, Grand Est, France (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

Amsterdam, Rotterdam, London, Manchester, Berlin, Hamburg, Budapest, Zagreb, Skopje, hybrid

DEPT® is a Growth Invention company built to help the world's most ambitious brands grow faster. Operating at the intersection of technology and marketing, our 4,000+ specialists deliver growth invention services across Brand & Media, Experience, Commerce, CRM, and Technology & Data. We're 50|50 tech and marketing, partner-led, and first to move. Clients include Google, Lufthansa, Meta, eBay, and OpenAI. We have been certified B Corp and Climate Neutral since 2021.

Job Purpose

We are seeking a talented and passionate AI Engineer to join our growing team. In this role, you will design, develop, and deploy AI solutions for our clients across various industries. You will collaborate with multidisciplinary teams to solve complex business challenges using cutting-edge data, machine learning and AI technologies.

What You'll Do

Design, develop and implement data-driven solutions tailored to client challenges, leveraging predictive models and AI technologies to deliver business value.Consult and guide clients by defining AI use cases, validating ideas through proof of concepts (POCs) and deploying scalable production-ready systems.Build and maintain end-to-end AI systems, with focus on scalability, observability and continuous optimization of model performance using MLOps principles.Develop multi-agent and AI workflows, built from scratch or using low-code tools, that reliably orchestrate…

https://www.linkedin.com/jobs/view/4402994122

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4402994122 · captured: 2026-04-18 · fit_reason: keyword + remote match

Lead Agentic AI Engineer [Riyadh / Dubai]

60% fit

Talent Seed · Greater Paris Metropolitan Region (On-site) · On-site · linkedin_jobs

ROLE

This role requires relocation to Riyadh/Dubai
You'll shape the technical direction and bring ambitious AI concepts to life, turning prototypes into scalable, reliable products that power next-generation customer experiences. Working closely with engineers, tech leads, and the CTO, you'll design systems that combine cutting-edge AI, strong infrastructure, and human-centred design.
Experienced Required:
AI EngineeringHands-on experience building complex, production-grade LLM-based systems across real-world use cases.Advanced prompt engineering skills, with a strong intuition for how models interpret, prioritise, and follow instructions in multi-step and long-context workflows.Deep understanding of LLM evaluation—recognising the limits of off-the-shelf metrics and able to design custom evaluation frameworks that measure what truly matters.Experience designing agentic systems, including task decomposition, tool creation and orchestration, and reliability engineering for non-deterministic workflows.Actively engaged with the evolving AI ecosystem—tracking research, experimenting with new capabilities, and translating emerging ideas into practical solutions.Comfortable going beyond established patterns to design novel approaches when existing techniques are insufficient.
Software EngineeringStrong grasp of software engineering fundamentals, including clean abstractions, separation of concerns, and writing maintainable, extensible code.Solid architectural judgment, with the abili…

https://www.linkedin.com/jobs/view/4353743722

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4353743722 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior AI/ML Software Engineer, YouTube Knowledge

60% fit

Google · Paris, Île-de-France, France (On-site) · On-site · linkedin_jobs

ROLE

Google welcomes people with disabilities.

Minimum qualifications:

Bachelor’s degree or equivalent practical experience.5 years of experience with software development in one or more programming languages.3 years of experience testing, maintaining, or launching software products, and 1 year of experience with software design and architecture.3 years of experience with one or more of the following: Speech/audio (e.g., technology duplicating and responding to the human voice), reinforcement learning (e.g., sequential decision making), ML infrastructure, or specialization in another ML field.3 years of experience with ML infrastructure (e.g., model deployment, model evaluation, optimization, data processing, debugging).

Preferred qualifications:

Master's degree or PhD in Computer Science, or a related technical field.5 years of experience with data structures and algorithms.1 year of experience in a technical leadership role.Experience developing accessible technologies.

About The Job

Google's software engineers develop the next-generation technologies that change how billions of users connect, explore, and interact with information and one another. Our products need to handle information at massive scale, and extend well beyond web search. We're looking for engineers who bring fresh ideas from all areas, including information retrieval, distributed computing, large-scale system design, networking and data storage, security, artificial intelligence, natural language process…

https://www.linkedin.com/jobs/view/4400783517

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400783517 · captured: 2026-04-18 · fit_reason: keyword + remote match

Applied AI Engineer

Sia · Paris, Île-de-France, France (On-site) · On-site · linkedin_jobs

60% fit

ROLE

Sia is a next-generation, global management consulting group. Founded in 1999, we were born digital. Today our strategy and management capabilities are augmented by data science, enhanced by creativity and driven by responsibility. We're optimists for change and we help clients initiate, navigate and benefit from transformation. We believe optimism is a force multiplier, helping clients to mitigate downside and maximize opportunity. With expertise across a broad range of sectors and services, our 3,000 consultants serve clients worldwide from 48 locations in 19 countries. Our expertise delivers results. Our optimism transforms outcomes.

Sia's AI & Data Business Unit is the powerhouse of our firm's innovation—merging cutting-edge Data Science, Generative AI, and advanced digital solutions to transform industries. With over 350 experts worldwide, we tackle projects from proof-of-concept to large-scale deployment, always pushing the boundaries of AI capabilities. Our 12 R&D labs in Europe and North America drive continuous research in areas like computer vision, MLOps, and deep learning, partnering closely with our business consultants for real-world impact. By joining Sia's AI & Data team, you'll step into a vibrant, collaborative environment that nurtures professional growth and empowers you to shape the future of AI-driven consulting.

Job Description

We are seeking a skilled Generative AI Engineer to join our team where you will harness model capabilities to implement …

https://www.linkedin.com/jobs/view/4402275155

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4402275155 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior AI / MLOps Engineer AWS - H/F

60% fit

Devoteam | AWS Partner · Toulouse, Occitanie, France (Hybrid) · On-site · linkedin_jobs

ROLE

Ø>Ý¸ &Bþ Présentation de l’équipe

Devoteam Toulouse, avec ses 200 collaborateurs, est un acteur incontournable de la transformation digitale sur le bassin toulousain. Au sein de l'agence, A Cloud, Partenaire de l'année 2025 d'AWS, a pour ambition d’aider ses clients à transformer la valeur de leurs données en succès durable et d’accélérer leur adoption de l’IA au sein de leurs équipes métiers.

Nous tirons parti à la fois de notre expertise sur les plateformes technologiques les plus innovantes et de notre connaissance métier avec la meilleure combinaison d’experts pour accompagner nos clients dans leur quête d’agilité, de compétitivité et de croissance durable.

Pour contribuer à ces évolutions, nous sommes à la recherche d’un(e) consultant(e) Senior MLOps Engineer AWS souhaitant rejoindre un collectif d’une dizaine de consultant(e)s spécialisé(e)s en data/ AI engineering sur AWS à Toulouse.

Description Du Poste

Ø=Üª Vos missions :

Concevoir et mettre en œuvre des architectures MLOps robustes et scalables sur AWSIndustrialiser le cycle de vie des modèles (training, validation, packaging, déploiement, monitoring, retraining)Mettre en place des pipelines CI/CD pour les workflows ML (Infrastructure as Code + ML pipelines)Automatiser les processus de déploiement et d’intégration des modèles en productionImplémenter des stratégies de monitoring des modèles (data drift, model drift, performance, observabilité)Collaborer étroitement avec les équipes Data Science, Data Engineering e…

https://www.linkedin.com/jobs/view/4390713132

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4390713132 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior GenAI Engineer (LLMs, RAG)

Recruit Action inc. · Toronto, ON (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

Senior Gen AI Engineer (LLMs, RAG)
Build next-generation AI solutions in the insurance industry using LLMs, RAG pipelines, vector databases, and Azure cloud. This hybrid Toronto role focuses on designing scalable GenAI systems, deploying production-ready AI services, and partnering with engineering and business teams to transform complex data into intelligent advisor tools.
What is in it for you:Salaried: $50-60 per hour.Incorporated Business Rate: $60-70 per hour.8-month contract with the potential for permanent employment.Full-time position: 37.50 hours per week.Remote on Monday and Friday; on-site Tuesday to Thursday.
Responsibilities: Architect and develop LLM-based solutions including retrieval-augmented generation (RAG) pipelines, embeddings, model fine-tuning, and evaluation frameworks.Build scalable Generative AI microservices and integrate them with internal enterprise systems.Perform advanced prompt engineering, agent design, and implement safety guardrails for AI systems.Evaluate open-source and commercial language models based on performance, cost, and risk.Collaborate with data teams to prepare training datasets, knowledge bases, and analytics pipelines.Manage ingestion and refresh processes for knowledge bases supporting RAG architectures.Implement monitoring and feedback loops to continuously improve model performance and solution quality.Partner with business stakeholders to define problem statements, data requirements, and delivery approaches.Document solutio…

https://www.linkedin.com/jobs/view/4402113555

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4402113555 · captured: 2026-04-18 · fit_reason: keyword + remote match

Forward Deployed AI Engineer, Canada - BCG X

60% fit

BCG X · Montreal, QC (Hybrid) · On-site · linkedin_jobs

ROLE

Locations: Toronto | Montreal

Who We Are

Boston Consulting Group partners with leaders in business and society to tackle their most important challenges and capture their greatest opportunities. BCG was the pioneer in business strategy when it was founded in 1963. Today, we help clients with total transformation-inspiring complex change, enabling organizations to grow, building competitive advantage, and driving bottom-line impact.

To succeed, organizations must blend digital and human capabilities. Our diverse, global teams bring deep industry and functional expertise and a range of perspectives to spark change. BCG delivers solutions through leading-edge management consulting along with technology and design, corporate and digital ventures—and business purpose. We work in a uniquely collaborative model across the firm and throughout all levels of the client organization, generating results that allow our clients to thrive.

Qui nous sommes

BCG X est l’unité de conception et de construction technologique de Boston Consulting Group (BCG). En tirant parti de la profonde expertise sectorielle et fonctionnelle de Boston Consulting Group (BCG), BCG X combine des connaissances technologiques avancées et un esprit entrepreneurial ambitieux pour aider les organisations à innover à grande échelle. Avec près de 3 /000 /technologues, scientifiques, programmeurs, ingénieurs et concepteurs centrés sur l’humain répartis dans plus de 80 /villes, BCG X conçoit et développe des plateformes …

https://www.linkedin.com/jobs/view/4389828108

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4389828108 · captured: 2026-04-18 · fit_reason: keyword + remote match

Staff AI/ML Engineer

Tetrix · Montreal, QC (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

'(Overview

Tetrix is a fast-growing startup on a mission to empower institutional investors to confidently invest in alternative markets (think Bloomberg, but for private markets). Our AI-native product transforms messy, unstructured private market data into real-time insights, combining automated document retrieval, AI-powered data extraction, and on-demand portfolio analytics - all in one seamless platform. We are growing quickly to becoming a category-defining product, and are working with some of the world's leading capital allocators. This is a unique opportunity to join a talented and driven team that wants to build next-generation infrastructure in financial markets, and is backed by stellar investors that include Innovation Endeavors, Eric Schmidt's VC firm, and angels from BlackRock, CPPIB, SoftBank, Lending Club, and Plaid.

Ø>Ý7Ø<ßýþ Who We Are

We are a small but mighty startup with deep industry background and expertise in both financial markets and modern technologies. We've worked hard to build a culture of trust, high ownership, and curiosity—where everyone feels empowered to contribute, challenge ideas, and grow together. We are led by our two co-founders, Olivier (CEO) & Naunidh (CTO), who met on the first day of school whilst pursuing their MBAs at Stanford and co-founded Tetrix in the Summer of 23.

Ø>Ý Ø<ßý Who You Are

You're proactive and hands-on – someone who takes initiative, figures things out, and keeps things moving. You learn and adapt quickly – someone …

https://www.linkedin.com/jobs/view/4400126518

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400126518 · captured: 2026-04-18 · fit_reason: keyword + remote match

Développeur.se principal.e, Applications IA (Sécurité) / Senior AI Dev (Safety)

60% fit

Mila - Quebec Artificial Intelligence Institute · Montreal, QC (Hybrid) · On-site · linkedin_jobs

ROLE

À propos de Mila

Fondé par le professeur Yoshua Bengio de l'Université de Montréal, Mila rassemble des chercheurs spécialisés en intelligence artificielle et plus précisément en apprentissage automatique, apprentissage profond et apprentissage par renforcement. Reconnu mondialement pour ses importantes contributions au domaine de l'apprentissage profond, Mila s'est particulièrement distingué dans la modélisation du langage, la traduction automatique, la reconnaissance d'objets et les modèles génératifs. Depuis 2017, Mila est le fruit d'une collaboration entre l'Université de Montréal et l'Université McGill, en lien étroit avec Polytechnique Montréal et HEC Montréal.

Mila s'est donné pour mission d'être un pôle mondial d'avancées scientifiques qui inspire l'innovation et l'essor de l'intelligence artificielle (IA) au bénéfice de tous.

Pour en connaitre davantage, veuillez consulter https://mila.quebec/

Description Du Poste

L'équipe Développement d'applications IA (DAI) conçoit et bâtit des applications alimentées par l'IA générative et l'apprentissage automatique. Nous travaillons étroitement avec les Studios, l'AMLRT et les équipes internes pour transformer des modèles, prototypes et idées en produits fonctionnels et démonstrations interactives.

En tant que développeur.se principal.e, vous serez affecté.e au Studio Sécurité de l'IA, où vous contribuerez à la conception et au développement d'une plateforme d'évaluation open source mesurant la sécurité et la robustesse co…

https://www.linkedin.com/jobs/view/4403119436

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4403119436 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI Engineer

Synechron · Toronto, ON (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

We areAt Synechron, we believe in the power of digital to transform businesses for the better. Our global consulting firm combines creativity and innovative technology to deliver industry-leading digital solutions. Synechron's progressive technologies and optimization strategies span end-to-end Artificial Intelligence, Consulting, Digital, Cloud & DevOps, Data, and Software Engineering, servicing an array of noteworthy financial services and technology firms. Through research and development initiatives in our FinLabs we develop solutions for modernization, from Artificial Intelligence and Blockchain to Data Science models, Digital Underwriting, mobile-first applications and more. Over the last 20+ years, our company has been honored with multiple employer awards, recognizing our commitment to our talented teams. With top clients to boast about, Synechron has a global workforce of 16,850+, and has 60 offices in 20 countries within key global markets.
Our challengeWe are seeking an AI Engineer with deep expertise in Generative AI, neural networks, and transfer learning to support a strategic data analytics platform. The candidate will lead the design and development of cutting-edge ML models and architectures involving zero-shot/few-shot learning and retrieval-augmented generation (RAG)
Additional Information*The base salary for this position will vary based on geography and other factors.In accordance with law, the base salary for this role if filled within Toronto, ON is C$1…

https://www.linkedin.com/jobs/view/4399748955

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399748955 · captured: 2026-04-18 · fit_reason: keyword + remote match

AI Developer – Agentic System

Xanadu · Toronto, ON (On-site) · On-site · linkedin_jobs

60% fit

ROLE

About Xanadu

Xanadu's mission is to build quantum computers that are useful and available to people everywhere.

At Xanadu, we are learners, innovators, researchers, collaborators and problem solvers. We are creating something that has never been built before. Few people in their life will be able to be a part of something like this, where if we are successful, the technologies we develop will solve some of the world's most challenging problems, and literally change the world. And that is something to be excited about!

Your Role And Responsibilities

As an AI Developer at Xanadu, you will play a crucial role in integrating generative AI into our R&D methodologies. This unique position allows you to contribute to the development of our internal tool stacks, aimed at advancing the first commercially viable quantum computer. The AI team's focus is to build and improve the modeling, optimization, simulation, data processing, and design methodology of our internal research. Situated at the intersection of multiple technical disciplines, you will have the opportunity to work with world leading researchers, scientists, engineers, and software developers. By infusing cutting-edge AI into software tools, you have the chance to super-charge their research progress and even reshape how they work. You will:

Design, develop, and maintain our agentic AI research system for accelerating internal R&D.Build automated processes and pipelines for structuring unstructured documents and execut…

https://www.linkedin.com/jobs/view/4402075034

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4402075034 · captured: 2026-04-18 · fit_reason: keyword + remote match

Generative AI Lead Developer

EXL · Greater Toronto Area, Canada (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

Job SummaryWe at EXL are seeking a highly skilled and passionate Generative AI Developer to join our growing team. This role demands deep expertise in Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), and prompt engineering, with a strong focus on driving business impact in the Retail and Consumer Packaged Goods (CPG) domain. The ideal candidate will be comfortable working at the intersection of technology and industry, building scalable and innovative AI solutions that support personalization, automation, and decision intelligence. The role also involves end-to-end project management, comprehensive documentation, proactive client and stakeholder communication, driving business development and account growth initiatives, along with strong technical expertise.
Key ResponsibilitiesDesign, develop, and deploy GenAI solutions using LLMs such as GPT, Claude, LLaMA, or open-source equivalents.Build and optimize RAG pipelines leveraging vector databases (e.g., FAISS, Pinecone, Weaviate) and embedding models.Engineer, evaluate, and refine prompts to improve output accuracy and relevance across different use cases.Collaborate with data scientists, engineers, and retail/CPG domain experts to identify high-impact AI opportunities.Monitor, maintain, and fine-tune deployed models to ensure performance, scalability, and business alignment.Stay current on advancements in the GenAI and LLM ecosystem and assess applicability to ongoing and future initiatives.
Required Qualif…

https://www.linkedin.com/jobs/view/4400790424

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400790424 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior Platform Engineer

60% fit

Mux · San Francisco, CA, US · On-site · ycombinator

Ø=Ü° $210K - $240K

ROLE

Senior Platform Engineer at Mux | Y Combinator Mux Video infrastructure for developers Senior Platform Engineer $210K - $240K • San Francisco, CA, US Job type Full-time Role Engineering, Devops Experience 6+ years Visa US citizen/visa only Skills Go, Kubernetes Apply to Mux and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Mux Mux is video for developers. Our mission is to democratize video by solving the hard problems developers face when building video: video encoding and streaming (Mux Video), video monitoring (Mux Data), and more. Video is a huge part of people's lives, and we want to help make it better. We're committed to building a healthy team that welcomes diverse backgrounds and experiences. We want people who care about our mission, are ready to grow, believe in our values (from Be Human to Turn Customers Into Fans), and want to improve the people around them. You'll join a tight-knit team with experience at places like Google, YouTube, Twitch, Reddit, Zencoder, Fastly, and more. Our founders previously started (and sold) Zencoder, an early leader in cloud video technology, and authored Video.js, the biggest HTML5 video player on the web. We organize Demuxed , the premier conference for video engineers in the world. We're backed by top investors like Coatue, Accel, Andreessen Horowitz, and Y Combinator. You'll get to work with amazing companies: hundreds of startups, plus Strava, Patreon, Vimeo, Robinhood, PB…

https://www.ycombinator.com/companies/mux/jobs/9A5uo2n-senior-platform-engineer

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:91159 · captured: 2026-04-18 · fit_reason: keyword + remote match

Engineering Manager, Growth

60% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States / Remote (US) · Remote ' · ycombinator

Ø=Ü° $197K - $274K

ROLE

Engineering Manager, Growth at Checkr | Y Combinator Checkr People infrastructure for the future of work Engineering Manager, Growth $197K - $274K • Denver, Colorado, United States; San Francisco, California, United States / Remote (US) Job type Full-time Role Engineering, Full stack Experience 11+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the Role Checkr is looking for an Engineering Manager to lead the Checkr Growth team. This team is hyper-focused on Checkr's self-serve business delivering products and capabilities for Checkr's growth funnel: Acquisition, Activation, Expansion & Retention. Leveraging your deep experience, you will build, coach, and inspire the engineering team to meet company objectives and bring new products and features to market. What You'll Do Lead a team to work on growth-focused ini…

https://www.ycombinator.com/companies/checkr/jobs/EW8zuA5-engineering-manager-growth

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:89051 · captured: 2026-04-18 · fit_reason: keyword + remote match

Founding Marketing Lead

60% fit

Resolve · New York, NY, US / San Francisco, CA, US / Boston, MA, US / Remote (US) · Remote · ycombinator

Ø=Ü° $90K - $150K

ROLE

Founding Marketing Lead at Resolve | Y Combinator Resolve Software to accelerate construction by finding issues in 3D BIM Founding Marketing Lead $90K - $150K • 0.01% - 0.50% • New York, NY, US / San Francisco, CA, US / Boston, MA, US / Remote (US) Job type Full-time Role Marketing Experience 3+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Russell Varriale Founder Russell Varriale Founder About the role About Resolve Resolve is web-based AI & VR review software that helps global construction teams build as fast as the future demands. The infrastructure software suite helps teams find hidden project risk before it derails schedules, drives rework, or drives up cost. Built for BIM-centric workflows, Resolve turns traditional reviews from a bottleneck into a competitive advantage — using immersive 3D environments, real-time collaboration, connected issue tracking, and intelligent spatial AI to align design, build, and operations teams. About the role We're making our first marketing hire at Resolve. This role is focused on creating, managing, and distributing content that drives real usage . We're looking for someone with 3–5+ years of product marketing or content experience who has actually made things: written content, recorded videos, shipped campaigns, and learned what works by doing. You'll be responsible for turning what we're building into clear, compelling content, getting it in front of the right people, …

https://www.ycombinator.com/companies/resolve/jobs/9ECcwEs-founding-marketing-lead

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92478 · captured: 2026-04-18 · fit_reason: keyword + remote match

Senior Sales Development Representative

60% fit

SnapMagic · Redwood City, CA, US / Remote (US) · Remote ' · ycombinator

Ø=Ü° $100 - $150

ROLE

Senior Sales Development Representative at SnapMagic | Y Combinator SnapMagic AI-assisted electronics design Senior Sales Development Representative $100 - $150 • 0.01% • Redwood City, CA, US / Remote (US) Job type Full-time Role Sales Experience 1+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Natasha Baker CEO Natasha Baker CEO About the role SnapMagic is building the digital growth engine for the global electronics industry. Our core platform (formerly SnapEDA) already reaches 2M+ engineers annually and is the default discovery layer for electronic components. We're now scaling SnapMagic Copilot, built on a proprietary dataset of 10M+ components, serving as the discovery and demand layer for the global electronics industry. We are hiring a Senior SDR to build, optimize, and scale SnapMagic's AI-driven outbound and top-of-funnel engine. You will: Architect AI-powered, signal-based outbound programs Implement AI agents to drive research, personalization, enrichment, and routing Own GTM tooling and automation systems end-to-end Improve conversion from first touch to qualified meeting Be responsible not just for booked meetings, but for measurable pipeline impact This role is ideal for someone who sees AI as leverage, not hype; and wants to apply it directly to revenue generation. What You'll Do Build an AI-Native GTM Engine Design and execute multi-channel outbound strategies powered by AI agents and signal-base…

https://www.ycombinator.com/companies/snapmagic/jobs/t7BgVaM-senior-sales-development-representative

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:90618 · captured: 2026-04-18 · fit_reason: keyword + remote match

Head of Sales

60% fit

SnapMagic · Redwood City, CA, US / Remote (San Francisco, CA, US) · Remote ' · ycombinator

Ø=Ü° $300K - $400K

ROLE

Head of Sales at SnapMagic | Y Combinator SnapMagic AI-assisted electronics design Head of Sales $300K - $400K • 0.02% - 0.50% • Redwood City, CA, US / Remote (San Francisco, CA, US) Job type Full-time Role Sales Experience 6+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Natasha Baker CEO Natasha Baker CEO About the role SnapMagic is building the digital growth engine for the global electronics industry. Our core platform (formerly SnapEDA) already reaches 2M+ engineers annually and is the default discovery layer for electronic components. We're now scaling SnapMagic Copilot, built on a proprietary dataset of 10M+ components, serving as the discovery and demand layer for the global electronics industry. This is a category-defining opportunity. We are specifically looking for someone who is deeply obsessed with the craft of selling – someone who studies it, practices it, and improves it continuously through direct customer conversations. Since our seed round, we've doubled revenue with a small team operating at a very high bar. The demand is real. The market is ready. What's missing is a sales leader who can turn momentum into a streamlined, repeatable engine. We're looking for someone who has personally built the sales motion at a VC-backed company and knows what it takes to scale from 7-figure to 8-figure revenue. What you'll do Read this carefully: this is a builder role , not a delegator role. You will perso…

https://www.ycombinator.com/companies/snapmagic/jobs/kXAprxJ-head-of-sales

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:86583 · captured: 2026-04-18 · fit_reason: keyword + remote match

Enterprise Account Executive

60% fit

SnapMagic · Redwood City, CA, US / Remote (San Francisco, CA, US) · Remote ' · ycombinator

Ø=Ü° $200K - $300K

ROLE

Enterprise Account Executive at SnapMagic | Y Combinator SnapMagic AI-assisted electronics design Enterprise Account Executive $200K - $300K • Redwood City, CA, US / Remote (San Francisco, CA, US) Job type Full-time Role Sales Experience 6+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Natasha Baker CEO Natasha Baker CEO About the role Please apply at our careers site: https://careers.snapmagic.com/o/account-executive-2-3 SnapMagic, the leading platform for electrical engineers, is revolutionizing electronics design with artificial intelligence. We help 1.5 million engineers design products like drones, electric cars, and virtual reality headsets 10x faster. With 10+ years of proven market traction (revenue-driven growth), SnapMagic is fresh off a round of VC funding from leading Silicon Valley investors to accelerate our business into the next stages of growth. What you'll do: As one of the foundational members of our Go To Market team, you'll be at the steering wheel of executing on our aggressive growth plans while simultaneously building the playbook for our growing team. You will own your sales process end to end, from pipeline generation, defining strategic sales plans, to forging strong partnerships with within our electronics ecosystem. What you'll be doing: Prospecting and Business Development: Define and execute on creative strategies to engage prospective clients through a variety of channels. Data Dr…

https://www.ycombinator.com/companies/snapmagic/jobs/GWuieQc-enterprise-account-executive

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:63247 · captured: 2026-04-18 · fit_reason: keyword + remote match

Director of Software & Product

60% fit

SnapMagic · Redwood City, CA, US / Remote (Redwood City, CA, US) · Remote ' · ycombinator

Ø=Ü° $200K - $250K

ROLE

Director of Software & Product at SnapMagic | Y Combinator SnapMagic AI-assisted electronics design Director of Software & Product $200K - $250K • 0.50% - 1.00% • Redwood City, CA, US / Remote (Redwood City, CA, US) Job type Full-time Role Engineering, Engineering manager Experience 6+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Natasha Baker CEO Natasha Baker CEO About the role Job description SnapMagic is building the digital growth engine for the global electronics industry. Our core platform (formerly SnapEDA) already reaches 2M+ engineers annually and is the default discovery layer for electronic components. We're now scaling SnapMagic Copilot, built on a proprietary dataset of 10M+ components, serving as the discovery and demand layer for the global electronics industry. This is a category-defining opportunity. We are looking for someone who is deeply obsessed with the craft of building products; someone who cares about building scalable systems, ships relentlessly, and improves through rapid feedback from real users. Since our seed round, we've doubled revenue with a small team operating at a very high bar. The demand is real. The market is ready. What's missing is a product and engineering leader who can turn momentum into a fast, repeatable shipping machine. We're looking for someone who has personally built and shipped complex systems at a fast-growing, VC-backed company and knows what it takes to sc…

https://www.ycombinator.com/companies/snapmagic/jobs/vZrseMD-director-of-software-product

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:86692 · captured: 2026-04-18 · fit_reason: keyword + remote match

Head of Operations

60% fit

SnapMagic · Redwood City, CA, US / Remote (Redwood City, CA, US) · Remote ' · ycombinator

Ø=Ü° $180K - $225K

ROLE

Head of Operations at SnapMagic | Y Combinator SnapMagic AI-assisted electronics design Head of Operations $180K - $225K • 0.20% - 1.00% • Redwood City, CA, US / Remote (Redwood City, CA, US) Job type Full-time Role Operations Experience 6+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Natasha Baker CEO Natasha Baker CEO About the role SnapMagic is building the digital growth engine for the global electronics industry. Our core platform, SnapMagic Search (formerly SnapEDA), already reaches 2M+ engineers annually and is the default discovery layer for electronic components. We're now scaling SnapMagic Copilot, built on a proprietary dataset of 10M+ components, serving as the discovery and demand layer for the global electronics industry. We've doubled revenue since our seed round with a small team operating at a very high bar. What we need now is someone to turn momentum into a machine. What you'll do We're hiring a Head of Operations to own and run all operations end-to-end. This is a hands-on builder role with full ownership of how the company operates day to day. Your mission is simple: Build and run a well-oiled operational machine so the company can scale at an ever-faster pace. You own clarity, cadence, and follow-through across the company. When something is unclear, broken, or slow, you take ownership of fixing the system. You will own business operations, people operations, internal tooling, and executiv…

https://www.ycombinator.com/companies/snapmagic/jobs/0XkjQVQ-head-of-operations

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:86691 · captured: 2026-04-18 · fit_reason: keyword + remote match

Sales Development Representative

60% fit

Deepgram · Remote / Remote (US) · Remote ' · ycombinator

Ø=Ü° $50K - $60K

ROLE

Sales Development Representative at Deepgram | Y Combinator Deepgram Building foundational AI for speech transcription and understanding. Sales Development Representative $50K - $60K • Remote / Remote (US) Job type Full-time Role Sales Experience 1+ years Visa US citizen/visa only Apply to Deepgram and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role Company Overview Deepgram is the leading platform underpinning the emerging trillion-dollar Voice AI economy, providing real-time APIs for speech-to-text (STT), text-to-speech (TTS), and building production-grade voice agents at scale. More than 200,000 developers and 1,300+ organizations build voice offerings that are ‘Powered by Deepgram’, including Twilio, Cloudflare, Sierra, Decagon, Vapi, Daily, Cresta, Granola, and Jack in the Box. Deepgram’s voice-native foundation models are accessed through cloud APIs or as self-hosted and on-premises software, with unmatched accuracy, low latency, and cost efficiency. Backed by a recent Series C led by leading global investors and strategic partners, Deepgram has processed over 50,000 years of audio and transcribed more than 1 trillion words. There is no organization in the world that understands voice better than Deepgram. Company Operating Rhythm At Deepgram, we expect an AI-first mindset—AI use and comfort aren’t optional, they’re core to how we operate, innovate, and measure performance. Every team member who works at Deepgram is expec…

https://www.ycombinator.com/companies/deepgram/jobs/eHsOFzM-sales-development-representative

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:78762 · captured: 2026-04-18 · fit_reason: keyword + remote match

Research Scientist, Voice

60% fit

Deepgram · US / Remote (US) · Remote ' · ycombinator

Ø=Ü° $150K - $250K

ROLE

Research Scientist, Voice at Deepgram | Y Combinator Deepgram Building foundational AI for speech transcription and understanding. Research Scientist, Voice $150K - $250K • US / Remote (US) Job type Full-time Role Science, Research Experience 3+ years Visa US citizen/visa only Skills Python Apply to Deepgram and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role Company Overview Deepgram is a foundational AI company building state of the art, production-ready AI models that streamline human-computer interaction and amplify productivity. By enabling seamless communication between humans and machines, we believe we can harness the untapped potential of AI and help pave the way for a more productive future. We passionately believe in the potential of audio data to transform lives, businesses, and interactions across the globe - which is why Deepgram is trusted by well-respected companies like NASA, Twilio, Auth0, and Spotify to push the boundaries of what is possible in voice technology! The Opportunity At Deepgram, we spend every day tackling big, real-world challenges in voice. Our customers hire us to solve their hardest problems, taking real, complex audio and transforming it into novel insights. And to raise the bar, everything we build needs scale in its DNA. We aren't content with simple horizontal scaling: we intend to replace entire data centers dedicated to speech analytics with a single rack of servers. These challenges de…

https://www.ycombinator.com/companies/deepgram/jobs/0wGOyli-research-scientist-voice

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:63170 · captured: 2026-04-18 · fit_reason: keyword + remote match

Corporate Account Executive

60% fit

Deepgram · Remote / Remote (US) · Remote ' · ycombinator

Ø=Ü° $240K - $260K

ROLE

Corporate Account Executive at Deepgram | Y Combinator Deepgram Building foundational AI for speech transcription and understanding. Corporate Account Executive $240K - $260K • Remote / Remote (US) Job type Full-time Role Operations Experience 6+ years Visa US citizen/visa only Apply to Deepgram and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role Company Overview Deepgram is the leading platform underpinning the emerging trillion-dollar Voice AI economy, providing real-time APIs for speech-to-text (STT), text-to-speech (TTS), and building production-grade voice agents at scale. More than 200,000 developers and 1,300+ organizations build voice offerings that are 'Powered by Deepgram', including Twilio, Cloudflare, Sierra, Decagon, Vapi, Daily, Cresta, Granola, and Jack in the Box. Deepgram's voice-native foundation models are accessed through cloud APIs or as self-hosted and on-premises software, with unmatched accuracy, low latency, and cost efficiency. Backed by a recent Series C led by leading global investors and strategic partners, Deepgram has processed over 50,000 years of audio and transcribed more than 1 trillion words. There is no organization in the world that understands voice better than Deepgram. Company Operating Rhythm At Deepgram, we expect an AI-first mindset—AI use and comfort aren't optional, they're core to how we operate, innovate, and measure performance. Every team member who works at Deepgram is expected…

https://www.ycombinator.com/companies/deepgram/jobs/V9YiHIJ-corporate-account-executive

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:81249 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer

60% fit

Connells Group · Milton Keynes, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

We are looking for a Platform Engineer to join our Group Technology team in Milton Keynes. You will be responsible for developing and maintaining components of the ConnellsX Platform. Your success is directly tied to enabling engineering teams to efficiently deploy, monitor, and scale their services on the platform. Your technical expertise and collaborative approach will drive both tactical implementations and contribute to strategic platform evolution.

You will be a hands-on engineer with a customer-centric mindset, working directly with development teams to understand their needs and deliver solutions. You will work across multiple technical domains including orchestration, automation, CI/CD pipelines, cloud services, observability, and security, developing deeper expertise in areas that align with platform priorities and your interests. Experience with Microsoft Azure is essential.

Daily coding is a core expectation of this role. You will implement platform features, automation tools, and infrastructure improvements while collaborating with team members to share knowledge and best practices. You will help establish Platform Engineering practices at Connells by mentoring colleagues and contributing to technical standards and documentation.

The platform team operates with a focus on automation and efficiency, maintaining a lean structure where engineers contribute across multiple areas of expertise. You will work alongside specialists in site reliability, onboarding, archit…

https://www.linkedin.com/jobs/view/4393563858

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4393563858 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer

60% fit

Cellular Origins · Cambridge, England, United Kingdom (On-site) · On-site · linkedin_jobs

ROLE

Company Description

Cell therapies have been shown to cure devastating and previously untreatable diseases like cancer by re-programming the patient's own immune system. Giving more patients access to these life-saving therapies requires a revolution in manufacturing and process automation.

Cellular Origins, is enabling scalable, cost-effective and efficient manufacture of cell and gene therapies. Our technology addresses the challenges associated with manufacturing at scale with an elegant robotic manufacturing automation solution that reduces cost and labour, eliminates human error, and increases efficiency.

Job Description

As a Platform Engineer at Cellular Origins, you'll work alongside Senior Engineers to design, build, and maintain the network and on-prem server infrastructure that powers our cell-therapy automation factory services. Partnering closely with our software-engineering teams, you'll help automate build tasks, create self-service environments, and promote smooth collaboration between dev and ops. In this role, you'll contribute to maintaining and extending the factory architecture—from rack-level physical layout through to servers, switching, firewalls, and storage fabrics. You'll assist with managing our on-prem hypervisors, IaaS/PaaS layers, and containerised workloads, gaining exposure to the full stack that supports our production systems.

Working within a DevOpsSec mindset, you'll help craft secure, repeatable deployment patterns for both the facto…

https://www.linkedin.com/jobs/view/4383241778

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4383241778 · captured: 2026-04-18 · fit_reason: keyword + remote match

Datacenter Deployment Engineer

60% fit

Nscale · London, England, United Kingdom (Remote) · Remote ' · linkedin_jobs

ROLE

About Nscale

Nscale is the GPU cloud engineered for AI. We provide cost-effective, high-performance infrastructure for AI start-ups and large enterprise customers. Nscale enables AI-focused companies to achieve superior results by reducing the complexity of AI development. Our GPU cloud bolsters technical capabilities and directly supports strategic business outcomes, including cost management, rapid innovation, and environmental responsibility.

At Nscale, our Deployment team plays a critical role in driving the delivery of our GPU infrastructure into our DCs.

We thrive on a culture of relentless innovation, ownership, and accountability, where every team member takes pride in their work and drives it with excellence and urgency. As an Nscaler, you'll build trust through openness and transparency, where everyone is inspired to do their best work. If you join our team, you'll be contributing to building the technology that powers the future

About The Role

We're looking for a Datacenter Deployment Engineer.

This role will be responsible for the execution of physical deployments, you will be engaged in the design and layout of large scale GPU infrastructure projects. This will range from ensuring BOMs are complete and compatible, through to laying out the physical hardware in racks, cabling, and everything inbetween.

What You'll Be Doing

Layout of IT equipment inside the datacentersTaking BOMs and ensuring all hardware fits into available space, applying best practice to…

https://www.linkedin.com/jobs/view/4392396186

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4392396186 · captured: 2026-04-18 · fit_reason: keyword + remote match

Cloud Operations & Infrastructure Graduate

60% fit

AVEVA · Cambridge, England, United Kingdom (Remote) · Remote ' · linkedin_jobs

ROLE

AVEVA is creating software trusted by over 90% of leading industrial companies.

AVEVA's Early Careers Recruiting team is actively searching for Software Development Graduates for 2026.

As a Cloud Developer Graduate, you will:

Work alongside our experienced cloud developers on one of our small, energetic development teams.

Learn the process of software development using modern languages like Agile to help full-time work in designing, coding, testing, documenting, and supporting our software products.

This is a brilliant opportunity to join the AVEVA R&D graduate program with opportunities across several disciplines including Engineering, Advanced Applications, Data Collection, Human-Machine Interfaces, and Industrial Platform.

You will build essential experience working with customer problems across diverse sectors such as marine, power, manufacturing and sustainability.

You will be supported to gain the relevant professional development to support your work and will be part of our AVEVA graduate development program. You will be part of the AVEVA graduate community and have the opportunity to take part in Action for Good Days and companywide business projects.

Advanced Apps

HMI/SCADA

PI

Industrial Platform

Engineering

Cloud Dev Ops

We are open to applications from a wide range of disciplines including Computer Science, Engineering, Mechanical Engineering.

UK Benefits include:

Flexible benefits fund, emergency leave days, adoption leave, 28 days annual leave (…

https://www.linkedin.com/jobs/view/4392878156

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4392878156 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer (Azure)

55 Exec Search · London Area, United Kingdom (Hybrid) · On-site · linkedin_jobs

60% fit

ROLE

We’re partnering with a boutique data analytics and technology consultancy with a strong and growing footprint across UK Central Government / Public Sector
Following sustained organic growth and increasing demand from major government departments, the consultancy is expanding its consulting team and seeking Azure-focused engineers with real Kubernetes experience who want to move further into consulting, solution design, and architecture.
We are open to candidates at different stages of their career: from Junior/mid Level Engineers ready to step into consulting, through to experienced Senior Consultants who can lead solution design and client conversations from day one.
This is not just a sys admin type of role and not a pure platform support position. It’s a delivery-focused consulting role, working closely with clients to design and implement secure, scalable, cloud-native solutions, with a clear pathway to grow into an Azure & Kubernetes Solution Architect.
Why Join Us?Join a small, fast-growing consultancy where your contribution genuinely mattersWork on long-term, meaningful programmes, including national-scale public sector projectsBe supported in your progression toward Azure & Kubernetes Solution ArchitectureBalance hands-on engineering with design, consulting, and technical leadershipAvoid rigid role definitions, your responsibilities will grow with your capability
What You’ll Be DoingDepending on your experience, you’ll be involved in:Designing and delivering Azure-…

https://www.linkedin.com/jobs/view/4357447980

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4357447980 · captured: 2026-04-18 · fit_reason: keyword + remote match

Infrastructure Engineer

Mercor · United Kingdom (Remote) · Remote ' · linkedin_jobs

60% fit

ROLE

About The Job

Mercor connects elite creative and technical talent with leading AI research labs. Headquartered in San Francisco, our investors include Benchmark, General Catalyst, Peter Thiel, Adam D'Angelo, Larry Summers, and Jack Dorsey.

Position: Software Engineer

Type: Freelance

Compensation: $75–$80/hour

Location: Remote

Duration: 1–2 weeks

Commitment: 15–40 hours/week

Role Responsibilities

Evaluate and compare the performance of AI-powered CLI coding agents on real-world infrastructure debugging tasks.Solve TerminalBench tasks using AI CLI agents to diagnose and fix failing systems.Write bash scripts to address root causes and ensure they survive service restarts.Diagnose broken systems inside Docker containers involving databases, TLS, pipelines, and more.Compare and rank agents' approaches after each task to inform product decisions.Work independently and asynchronously with flexible hours.

Qualifications

Must-Have

3+ years of experience in software engineering with hands-on debugging of systems and infrastructure.Strong bash/shell scripting proficiency.Experience with Docker and containerization.Skills in infrastructure and systems debugging with PostgreSQL, MySQL, Redis, nginx, TLS, systemd, and log analysis.Familiarity with version control workflows (Git, PRs, issue tracking).

Preferred

Experience with AI coding tools like Copilot, Cursor, or Claude.

Start Date

Immediate

Application Process (Takes 20–30 mins to complete)

Upload resumeAI interview:…
https://www.linkedin.com/jobs/view/4400367546

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400367546 · captured: 2026-04-18 · fit_reason: keyword + remote match

Platform Engineer

60% fit

Oakwell Hampton Group · London Area, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

A high-growth climate analytics company is looking for a Platform Engineer to help scale the infrastructure behind a global climate risk platform used by enterprise clients. This is an opportunity to work on large-scale geospatial datasets, high-performance compute workloads, and cloud-native infrastructure in a mission-driven environment.
You'll collaborate closely with platform, data, and geospatial engineering teams to build reliable, scalable systems that power advanced climate modelling and analytics.
What You'll Be DoingDesign and build scalable AWS cloud infrastructureManage infrastructure using Infrastructure as Code (Terraform, CloudFormation)Develop and improve CI/CD pipelines and deployment workflowsOptimise systems handling large geospatial datasets and compute workloadsImplement monitoring, logging, and alertingImprove platform security and reliabilityPartner with engineering teams to support delivery and scalingDrive cloud cost optimisation initiatives
Tech EnvironmentAWS (EC2, ECS/EKS, Lambda, S3, RDS, DynamoDB, VPC)Terraform / CloudFormationKubernetes & DockerCI/CD (GitHub Actions, GitLab, Jenkins)Python / Bash scriptingObservability (Prometheus, Grafana, Datadog)
Requirements3+ years experience in Platform / DevOps / Cloud EngineeringStrong hands-on AWS experienceInfrastructure as Code (Terraform or similar)Containerisation (Docker) and orchestration (Kubernetes/ECS)CI/CD pipeline development experienceMonitoring and observability toolingStrong scripting skil…

https://www.linkedin.com/jobs/view/4394610664

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4394610664 · captured: 2026-04-18 · fit_reason: keyword + remote match

Sr Cloud Platform Engineer WOTC

Wizards of the Coast · Montreal, QC (On-site) · On-site · linkedin_jobs

60% fit

ROLE

At Wizards of the Coast, we connect people around the world through play and imagination. From our genre-defining games like Magic: The Gathering® and Dungeons & Dragons® to our growing multiverse, we continue to innovate and build new ways to foster friendship and connection. That's where you come in!

We are seeking a Senior Cloud Platform Engineer to join our Central Technology Build Engineering team as we develop and curate an Unreal Engine ecosystem to share across our internal game studios. You will join a specialized team responsible for unblocking developers, ensuring build stability, and optimizing iteration workflows across multiple concurrent titles. This is a hands-on IC role where engineering excellence meets operational reliability. You'll collaborate closely with game development teams to ensure a seamless experience for all Unreal-powered studios under our banner. In this role, you will design, build, and run scalable AWS architectures that serve as the backbone for deploying our Unreal Engine build systems and tools.

What You'll Do:

Work with the team to design a ready-made cloud-first suite of Unreal Engine build systems and tooling Develop and maintain infrastructure-as-code to ensure consistency and repeatability Support ongoing AWS operations for build infrastructure Support and administer studio-specific Perforce environments Work with studio teams to optimize our systems for cost, scale, and latency Define standards and best practices for performance,…

https://www.linkedin.com/jobs/view/4383454254

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4383454254 · captured: 2026-04-18 · fit_reason: keyword + remote match

Python Engineer, Truework

55% fit

Checkr · San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $150K - $176K

ROLE

Python Engineer, Truework at Checkr | Y Combinator Checkr People infrastructure for the future of work Python Engineer, Truework $150K - $176K • San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 3+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the Role Truework, recently acquired by Checkr, provides access to crucial employment and income information needed for mortgage loans, apartment rentals, background checks and more. We are disrupting a $5B+ industry, dominated by legacy incumbents, with a better product and intense customer focus. Within Checkr, Truework operates independently, allowing rapid iteration while also being able to leverage the strengths and assets from our core business. Think of it like joining a startup within a startup! As a Software Engine…

https://www.ycombinator.com/companies/checkr/jobs/vXk4adq-python-engineer-truework

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92278 · captured: 2026-04-18 · fit_reason: low fit

Senior Python Engineer, Truework

55% fit

Checkr · San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $185K - $218K

ROLE

Senior Python Engineer, Truework at Checkr | Y Combinator Checkr People infrastructure for the future of work Senior Python Engineer, Truework $185K - $218K • San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the Role Truework, recently acquired by Checkr, provides access to crucial employment and income information needed for mortgage loans, apartment rentals, background checks, and more. We are disrupting a $5B+ industry, dominated by legacy incumbents, with a better product and intense customer focus. Within Checkr, Truework operates independently, allowing rapid iteration while also being able to leverage the strengths and assets from our core business. Think of it like joining a startup within a startup! As a …

https://www.ycombinator.com/companies/checkr/jobs/XxaPwjN-senior-python-engineer-truework

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92286 · captured: 2026-04-18 · fit_reason: low fit

Enterprise AI Strategist, Deepgram for Restaurants

55% fit

Deepgram · San Francisco, CA · On-site · ycombinator

ROLE

Enterprise AI Strategist, Deepgram for Restaurants at Deepgram | Y Combinator Deepgram Building foundational AI for speech transcription and understanding. Enterprise AI Strategist, Deepgram for Restaurants • San Francisco, CA Job type Full-time Role Operations Visa US citizen/visa only Apply to Deepgram and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role Company Overview Deepgram is the leading platform underpinning the emerging trillion-dollar Voice AI economy, providing real-time APIs for speech-to-text (STT), text-to-speech (TTS), and building production-grade voice agents at scale. More than 200,000 developers and 1,300+ organizations build voice offerings that are 'Powered by Deepgram', including Twilio, Cloudflare, Sierra, Decagon, Vapi, Daily, Cresta, Granola, and Jack in the Box. Deepgram's voice-native foundation models are accessed through cloud APIs or as self-hosted and on-premises software, with unmatched accuracy, low latency, and cost efficiency. Backed by a recent Series C led by leading global investors and strategic partners, Deepgram has processed over 50,000 years of audio and transcribed more than 1 trillion words. There is no organization in the world that understands voice better than Deepgram. Company Operating Rhythm At Deepgram, we expect an AI-first mindset—AI use and comfort aren't optional, they're core to how we operate, innovate, and measure performance. Every team member who works at Deepgram is…

https://www.ycombinator.com/companies/deepgram/jobs/rzjRzjp-enterprise-ai-strategist-deepgram-for-restaurants

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:88498 · captured: 2026-04-18 · fit_reason: low fit

Staff AI Engineer - Utah

55% fit

Podium · Lehi, UT, US · On-site · ycombinator

Ø=Ü° $190K - $215K

ROLE

Staff AI Engineer - Utah at Podium | Y Combinator Podium Get more leads. Make more money. Staff AI Engineer - Utah $190K - $215K • Lehi, UT, US Job type Full-time Role Engineering, Full stack Experience 11+ years Visa US citizen/visa only Apply to Podium and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role At Podium, our mission is to arm every local business with a complete platform and outcome-driven AI employees that convert leads into real, paying customers. Every day, millions of workers use our AI lead conversion and communication platform to help them get more leads and make more money. Our work and focus on helping local businesses thrive has been recognized across the industry, including Forbes' Next Billion Dollar Startups, Forbes' Cloud 100, the Inc. 5000, and Fast Company's World's Most Innovative Companies. At Podium, we believe in fostering a culture that thrives on hiring and developing exceptional talent. Our operating principles serve as a compass, guiding daily behavior and decision-making, and ensure we hire people who will thrive at Podium. If you resonate with our operating principles and are energized by our mission, Podium will be a great place for you! Staff AI Engineer The Role We are looking for a talented Staff AI Engineer to help build and scale our AI Agent platform—a powerful system that autonomously interacts with customers, handling millions of conversations every month. In this role, you'll desig…

https://www.ycombinator.com/companies/podium/jobs/l8BTq6f-staff-ai-engineer-utah

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:88045 · captured: 2026-04-18 · fit_reason: low fit

Senior AI Engineer - Utah

55% fit

Podium · Lehi, UT, US · On-site · ycombinator

Ø=Ü° $160K - $190K

ROLE

Senior AI Engineer - Utah at Podium | Y Combinator Podium Get more leads. Make more money. Senior AI Engineer - Utah $160K - $190K • Lehi, UT, US Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Podium and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role At Podium, our mission is to arm every local business with a complete platform and outcome-driven AI employees that convert leads into real, paying customers. Every day, millions of workers use our AI lead conversion and communication platform to help them get more leads and make more money. Our work and focus on helping local businesses thrive has been recognized across the industry, including Forbes' Next Billion Dollar Startups, Forbes' Cloud 100, the Inc. 5000, and Fast Company's World's Most Innovative Companies. At Podium, we believe in fostering a culture that thrives on hiring and developing exceptional talent. Our operating principles serve as a compass, guiding daily behavior and decision-making, and ensure we hire people who will thrive at Podium. If you resonate with our operating principles and are energized by our mission, Podium will be a great place for you! The Role Podium is looking for a talented Senior AI Engineer to help build and scale our AI Agent platform —a powerful system that autonomously interacts with customers, handling millions of conversations every month. In this role, you'll design and delive…

https://www.ycombinator.com/companies/podium/jobs/pp9fqBd-senior-ai-engineer-utah

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:88044 · captured: 2026-04-18 · fit_reason: low fit

Fullstack Software Engineer (x/f/m) - AI Customer Support

55% fit

Alan · Montpellier, Occitanie, France (Hybrid) · On-site · linkedin_jobs

ROLE

You. Better. With Alan.

Alan is building a vertically integrated health partner that unites insurance and smart healthcare delivery into one seamless system. Our vision is to make prevention the new norm of care for all.

Our mission: Help people live in good health to 100 while helping employers feel proud, turning health benefits from a cost centre into their most valuable investment.

By connecting all aspects of care (private, public, and direct to consumer) we create the most member-centric healthcare experience, reducing claims costs while generating new monetization opportunities.

We partner with tens of thousands of companies across France Ø<ÝëØ<Ý÷, Spain Ø<ÝêØ<Ýø, Belgium Ø<ÝçØ<Ýê, and C
million members.

How we work: our Leadership Principles

Mission is the Boss — We think long-term and are methodical optimists who take risks, seeking our mission's success above all else.Distributed Ownership — Accountable enlightened despots: everyone owns their decisions and results.Radical Transparency — All information is accessible and written-first, so everyone can make the best decisions asynchronously.Always Growing — Direct, positive, and caring feedback, combined with self-growth ownership.

Ø=Ü< Who we are

You. Better. With Alan.

Alan's vision is to make prevention the new norm of care for all. Our mission is to help people live in good health to 100 while helping employers feel proud, turning health benefits from a cost center into their most valuable in…

https://www.linkedin.com/jobs/view/4376804352

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4376804352 · captured: 2026-04-18 · fit_reason: low fit

Staff AI Engineer - Notebooks

55% fit

Datadog · Paris, Île-de-France, France (Hybrid) · On-site · linkedin_jobs

ROLE

We're looking for a Staff Software Engineer with deep experience in AI/ML to join the Notebooks team. Datadog Notebooks provide a flexible tool for ad-hoc data analysis, collaborative investigations, and technical documentation. Its fluidity makes it an ideal surface for agents and humans to interact. The Notebook team is building AI-driven workflows for natural-language data analysis, assisted graph generation, automated investigations, and more. You will act as a technical leader in this domain, focused on building industry-leader AI-powered features. We're looking for a product-minded ML engineer with strong technical expertise, excellent communication skills, and a track record of driving impactful initiatives end to end.

At Datadog, we place value in our office culture - the relationships and collaboration it builds and the creativity it brings to the table. We operate as a hybrid workplace to ensure our Datadogs can create a work-life harmony that best fits them.

What You'll Do:

Spearhead AI initiatives within Notebooks, serving as the technical owner for the design, development, and deployment of ML/AI-driven features.Shape the long-term technical direction and strategy for AI workflows within Notebooks and related products.Champion the full development cycle: from rapid prototyping and defining success metrics/evaluations to running experiments, iterating, and achieving production-ready scale and reliability.Drive product direction with a strong user focus, always …

https://www.linkedin.com/jobs/view/4367317892

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4367317892 · captured: 2026-04-18 · fit_reason: low fit

Senior AI Agent Engineer

Mirakl · Bordeaux, Nouvelle-Aquitaine, France (On-site) · On-site · linkedin_jobs

55% fit

ROLE

About Mirakl:Founded in 2012, Mirakl has been at the forefront of marketplace innovation, empowering every business to compete in the platform economy.

Today, Mirakl's operating system combines an enterprise marketplace solution (Mirakl Platform) that enables retailers and B2B organizations to launch, scale, and operate marketplaces and dropship, AI-powered multichannel selling (Mirakl Connect), retail media (Mirakl Ads) and an agentic commerce infrastructure (Mirakl Nexus).

With dual headquarters in Boston and Paris, Mirakl helps a global ecosystem of 450+ marketplaces (B2C and B2B) and a network of over 100k third-party marketplace sellers. Brands like Macy's, Decathlon, Carrefour, Asos, and Airbus Helicopters use Mirakl to grow their businesses in new and remarkable ways.For more information: www.mirakl.com.Mirakl in Numbers:Ø=ÝÓþ Founded in 2012 | Member of French Tech Next40Ø=Üe 750+ employees in 9 offices worldwide: Paris, Barcelona, Bordeaux, Boston, London, Munich, New York, Sydney, TokyoØ<ÝëØ<Ý÷ 350+ Mirakl Tech teams members mainly based in France&™þ 5 Saas SolutionsOur Values:Working at M means accelerating your career alongside ambitious, passionate, and supportive colleagues. We're proud of the diversity of backgrounds, perspectives, and experiences that make our teams unique.Our 5 values guide how we collaborate:Ø=Ü¡ Work Hard Together: Teamwork and collaboration are the foundation of our successØ<ßÆ Get Things Done: We prioritize action and efficiency for impactful …

https://www.linkedin.com/jobs/view/4385723864

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4385723864 · captured: 2026-04-18 · fit_reason: low fit

AI Developer

Sparagus · Greater Dunkerque Area (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

Senior AI Developer

Ø=ÜÍDunkirk, France (Hybrid) | Industrial / Manufacturing Sector

We are supporting a major industrial player in deploying a multilingual AI Service Desk agent (7 languages).

As a Senior AI Developer, you will take AI prototypes into production and build the core components of a scalable AI system.

Role

Develop the AI conversation engine (LLM integration, structured outputs, fallback logic)Build the voice pipeline (STT/TTS, real-time translation)Implement RAG systems (embeddings, semantic search, knowledge indexing)Create automated QA frameworks for testing and evaluationContribute to scaling and new use cases

Profile

Strong experience in Python and Node.js / TypeScriptProven experience with LLMs in production (prompting, tool use, orchestration)Hands-on with RAG/vector databases / semantic searchExperience with speech technologies (Deepgram, Azure Speech, ElevenLabs)Familiar with testing frameworks (Pytest, Jest)Ø<ÝëØ<Ý÷ Fluent French + Ø<ÝìØ<Ýç English (B2+)

https://www.linkedin.com/jobs/view/4392990933

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4392990933 · captured: 2026-04-18 · fit_reason: low fit

LEAD AI ENGINEER

55% fit

LCL · Villejuif, Île-de-France, France (On-site) · On-site · linkedin_jobs

ROLE

Le contexte : Le pôle IA, rattachée à la Direction des systèmes d’information de LCL, est une équipe à taille humaine composée d’une équipe scientifique (Data Science & Data Engineering), trois chefs de projets scientifiques, un référent architecture et de son responsable. Il a vocation à proposer des solutions innovantes, basées sur la Data Science, pour répondre aux divers besoins métiers de la banque, dans une optique d’industrialisation des projets menés.
Le poste : En tant que Lead AI Engineer, vous jouerez un rôle central dans la définition de la stratégie technique, l'encadrement des ingénieurs et la garantie de l'excellence opérationnelle des solutions d'IA déployées. Vous serez le garant de la qualité, de la robustesse et de la scalabilité des projets, depuis la preuve de concept jusqu'à leur maintien en production à grande échelle.
Vos missions principales :Définir la vision technique et l'architecture pour l'industrialisation des projets d'IA (MLOps et GenAI Ops).Encadrer et mentorer l'équipe d'AI Engineers, en étant le référent technique sur les bonnes pratiques.Piloter l'automatisation (CI/CD) et l'optimisation de l'infrastructure de déploiement (Kubernetes, Cloud, GPU).Garantir la fiabilité, la sécurité et le monitoring des solutions d'IA en production.Diriger l'industrialisation des projets d'IA Générative (LLM), de la gestion des prompts au déploiement.Assurer une veille technologique et être le point de contact technique principal pour les autres équipes.

P…

https://www.linkedin.com/jobs/view/4396459016

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4396459016 · captured: 2026-04-18 · fit_reason: low fit

Tech Lead AI Engineer – Développement Chatbot (H/F) - CDI - Lille

55% fit

Logical Conseils · Lille, Hauts-de-France, France (Hybrid) · On-site · linkedin_jobs

ROLE

Intégré(e) à notre pôle d'expertise en intelligence artificielle, en tant que Lead AI Engineer vous interviendrez auprès de nos clients pour concevoir et déployer des chatbots reposant sur des technologies d'IA générative.Ces solutions, accessibles via les sites web et applications mobiles, auront pour vocation de délivrer des réponses fiables, contextualisées et en temps réel, tout en accompagnant les utilisateurs de manière proactive tout au long de leur parcours d'achat.
Vos missionsEn tant qu'expert / lead en IA générative, vous interviendrez sur :La conception et la mise en œuvre d'un agent conversationnelL'utilisation de frameworks avancés tels que LangChain et LangGraphL'intégration de bases vectorielles pour améliorer la pertinence des réponsesLe développement de pipelines IA robustes et industrialisablesL'accompagnement technique des équipes sur les bonnes pratiques IALa mise en production de solutions agentiques performantesL'évaluation et l'amélioration des modèles (tests, qualité, pertinence)
Livrables attendusAudit technique et recommandationsConseil en architecture IA générativeDéveloppement de solutions Data ScienceMise en production d'un agent conversationnel
Profil recherchéCompétences techniques :Python : confirméPySpark : notions / intermédiaireLangChain : avancéLangGraph : avancéExpérience sur les architectures RAG et bases vectoriellesMaîtrise des outils de tests d'agents (type Ragas)Connaissance des approches de modélisation (type MCP)
Expérience :Expéri…

https://www.linkedin.com/jobs/view/4394606146

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4394606146 · captured: 2026-04-18 · fit_reason: low fit

Ingénieur IA Generative sénior F/H

55% fit

Orange Business · Balma, Occitanie, France (Hybrid) · On-site · linkedin_jobs

ROLE

Votre rôle

Qui sommes-nous ?

Digital Services est l'ESN d'Orange Business qui combine les expertises Cloud, Digital, Data et IA. Plus de 4 000 experts répartis dans toute la France accompagnent des entreprises et institutions publiques dans leur transformation numérique. Nous rejoindre, c'est participer à l'accompagnement de nos clients à travers des missions stimulantes et variées.

Intégré au pôle Cloud Big Data & IA de notre Agence Occitanie basée à Balma, vous rejoindrez une équipe dynamique dédiée à des projets techniques chez nos clients de proximité dans les domaines du cloud, de l'intelligence artificielle et de l'analyse/traitement de données. En tant que Développeur/Ingénieur IA, vous jouerez un rôle clé dans la conception, l'optimisation et l'intégration de modèles d'IA.

Missions principales

Vous contribuerez au développement et à la mise en oeuvre de solution d'intelligence artificielle pour répondre aux besoins stratégiques et opérationnels des clients, et serez amené à accompagner les parties prenantes dans l'adoption de ces technologies.

Conseiller les clients sur l’adoption, l'application, la contextualisation et l'intégration de l'IA générative dans leurs processus métier Identifier, évaluer et définir les uses cases en IA générative pour les clients, en analysant leur apport business, leur complexité, la faisabilité et en animant des ateliers d’idéation. Développer des applications d’IA Générative : prompting, agent, orchestrateur, pipeline RAG, base …

https://www.linkedin.com/jobs/view/4361430710

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4361430710 · captured: 2026-04-18 · fit_reason: low fit

AI Engineer H/F

Akkodis · Bordeaux, Nouvelle-Aquitaine, France (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

Akkodis, est un acteur mondial de l'ingénierie et de l'IT et un leader dans la smart industrie. Nous accompagnons la transformation digitale de nos clients à travers quatre services : Consulting, Solutions, Talent et Academy. Nous co-créons et imaginons des solutions de pointe pour répondre aux défis majeurs de notre société. Informations sur l'entrepriseAkkodis, est un acteur mondial de l'ingénierie et de l'IT et un leader dans la smart industrie. Nous accompagnons la transformation digitale de nos clients à travers quatre services : Consulting, Solutions, Talent et Academy. Nous co-créons et imaginons des solutions de pointe pour répondre aux défis majeurs de notre société.Présents dans 30 pays, nos 50 000 experts en IT et en ingénierie allient expertise technologique et connaissance sectorielle pour bâtir un avenir durable, portés par une culture d'inclusion et de diversité.Akkodis en France, ce sont près de 9.000 experts répartis sur l'ensemble du territoire, des collaborateurs partageant des valeurs fortes d'honnêteté, de respect, d'équité et d'inclusion. Notre engagement : leur permettre au quotidien d'être eux-mêmes au travail, et acteurs de leur vie et de leur développement au sein d'Akkodis.Akkodis est une entreprise handi-engagée et inclusive. Tous nos postes sont ouverts aux handicaps et à la diversité. Tous différents, tous compétents ! Akkodis est la marque technologique de The Adecco GroupMissionsLa ligne de service Consulting & Solutions d'Akkodis France renfor…

https://www.linkedin.com/jobs/view/4390820994

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4390820994 · captured: 2026-04-18 · fit_reason: low fit

Senior Math Libraries Engineer – Emulation in AI and HPC

55% fit

NVIDIA · Marseille, Provence-Alpes-Côte d'Azur, France (Hybrid) · On-site · linkedin_jobs

ROLE

We are looking for software engineers to join our math libraries teams for AI and HPC kernel generation, specifically targeting emulation of math operations across different precisions. Around the world, leading commercial and academic organizations are revolutionizing AI, scientific and engineering simulations, and data analytics, using data centers powered by GPUs. Applications of these technologies are in healthcare, NLP, VR, deep learning, autonomous vehicles and countless others. Did you know our team develops the GPU accelerated math libraries that makes all of this possible? If the idea of tinkering with bits and precision formats in math operations and applying your knowledge to develop and optimize algorithms to make an impact around world excite you, come and join our team!

What You Will Be Doing

Scoping, designing, and implementing high quality and performance numerical dense linear algebra software on GPUs.Providing technical leadership and feedback to library engineers working with you on projects and sometimes mentor interns.Working closely with product management and other internal and external customers to understand feature and performance requirements and help define the technical roadmaps of libraries.Finding opportunities to improve library performance and reduce code maintenance overhead through re-architecting.

What We Need To See

PhD or Master's degree in Computer Science, Applied Math, or related science or engineering field of study (or equivalent…

https://www.linkedin.com/jobs/view/4395471735

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4395471735 · captured: 2026-04-18 · fit_reason: low fit

AI Engineer

Publicis Re:Sources · Paris, Île-de-France, France (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

Description De L'entreprise

Publicis Re:Sources is the global backbone of Publicis Groupe, delivering advanced technology and operational solutions. The GenAI team leads innovation in generative AI, developing and deploying large language models (LLMs) and advanced AI technologies for business impact.

Descriptif du poste

We are seeking an experienced AI/ML Engineer to design, build, and scale our generative AI solutions. In this role, you will lead initiatives around large language models, image and multimodal generation, and the deployment of these technologies in production using Google Cloud Platform (GCP). You will set technical standards for GenAI, mentor the team, and ensure best practices in software engineering, collaboration, and code quality.

Responsabilités

Architect and develop backend services, APIs, and orchestration layers that power GenAI features, including LLM-driven workflows (RAG, agents, content generation) and image generation pipelines.Own end-to-end delivery of GenAI-powered features: from system design and implementation to deployment, monitoring, and iteration.Build and maintain scalable, production-grade infrastructure across GCP and Azure, leveraging services such as Vertex AI, Cloud Run, Azure OpenAI Service, Azure Web Apps, and Azure Application Insights.Design and maintain CI/CD pipelines for GenAI applications, ensuring reliable deployment, versioning, and rollback of model integrations and service dependencies.Enforce software engineering…

https://www.linkedin.com/jobs/view/4398890100

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4398890100 · captured: 2026-04-18 · fit_reason: low fit

Staff AI Engineer

Gorgias · Paris, Île-de-France, France (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

We believe conversations will become the #1 way to shop.

At Gorgias, we’re building the platform that makes this real: a unified AI agent that sells, supports, and re-engages customers across the entire journey. Conversational Commerce is the future of ecommerce, and we’re leading that shift.

Our mission is to turn every interaction between a brand and its customers into a relationship: personal, seamless, and intelligent. By combining deep product expertise with the latest in AI, we’re making shopping feel more natural, human, and connected than ever before.

To win, we focus relentlessly on:

Quality: conversations that feel authentic and on-brand.Experience: effortless shopping from chat to checkout.Re-engagement: personal, 1-1 dialogue instead of noisy marketing.

The opportunity is massive. As AI reshapes how people buy, Gorgias is building the foundation for the next decade of ecommerce, where every brand has its own intelligent agent and every customer feels understood.

Join us to make Conversational Commerce real.

Team & Context

Gorgias is an AI-first company building products powered by LLMs and agent-based systems.

As we scale our AI capabilities, we need to improve how we evaluate, iterate, and operate these systems in production. Today, parts of this process remain manual or fragmented, especially around prompt iteration, validation, and evaluation workflows.

This role will focus on building and scaling the systems that support AI evaluation and iteration, …

https://www.linkedin.com/jobs/view/4397318696

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4397318696 · captured: 2026-04-18 · fit_reason: low fit

Ingénieur en science des données Intelligence Artificielle expérimenté H/F - IBM…

55% fit

IBM · Strasbourg, Grand Est, France (Hybrid) · On-site · linkedin_jobs

ROLE

Introduction

Une carrière chez IBM Consulting repose sur des relations clients durables et une étroite collaboration à l'échelle mondiale. Vous travaillerez avec des entreprises de premier plan issues de tous les secteurs, en les aidant à définir leur stratégie en matière de cloud hybride et d'intelligence artificielle. Grâce au soutien de nos partenaires stratégiques, à la technologie robuste d'IBM et à Red Hat, vous disposerez des outils nécessaires pour susciter des changements significatifs et accélérer l'impact sur nos clients. Chez IBM Consulting, la curiosité est le moteur de la réussite. Vous serez encouragé à remettre en question les normes, à explorer de nouvelles idées et à créer des solutions innovantes qui produisent des résultats concrets. Notre culture de croissance et d'empathie met l'accent sur votre développement de carrière à long terme tout en valorisant vos compétences et vos expériences uniques.

Your Role And Responsibilities

Dans un contexte d’accélération des usages de l’intelligence artificielle, nous recherchons un(e) Ingénieur en science des données Intelligence Artificielle expérimenté capable d’intervenir sur l’ensemble du cycle de vie d’un produit IA, depuis l’identification des opportunités métier jusqu’à l’industrialisation et au suivi en production. Ce rôle couvre aussi bien le machine learning classique que le deep learning, l’IA générative, l’IA agentique, les modèles multimodaux, l’évaluation de modèles, l’ingénierie de données pour l’IA…

https://www.linkedin.com/jobs/view/4400423389

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400423389 · captured: 2026-04-18 · fit_reason: low fit

Ingénieur/Ingénieure Intelligence Artificielle secteur nucléaire

55% fit

Up Technologies · Aix-en-Provence, Provence-Alpes-Côte d'Azur, France (On-site) · On-site · linkedin_jobs

ROLE

Up Technologies est une société de conseil pour l'industrie avec des agences sur Lyon, Grenoble, Aix-en-Provence, Toulon, Nice Sophia et Paris.
Nous sommes spécialisés dans l'électronique, des systèmes embarqués et le logiciel pour l'industrie. Nous intervenons en assistance technique ainsi qu'au forfait à travers notre bureau d'études pour fabriquer des cartes électroniques et briques logicielles.
Forts d'expériences dans des sociétés de conseil en tant que Managers et Consultants, notre ambition est de proposer une société qui associe les objectifs professionnels des Consultants et ceux d'une société de conseil en ingénierie à taille humaine.
Le point commun des Consultants d'Up Technologies est d'avoir un esprit de services orienté Client, une capacité d'autonomie et une aisance relationnelle. La réussite des missions est portée par l'écoute et la transparence du trinôme Consultant, Client, Manager.
Vous partagez ces principes ? Vous souhaitez vous investir et relever des challenges ? Alors rejoignez-nous et venez participer à notre développement !
Are you Up for it ?
Afin d'accompagner notre client spécialisé dans l'inspection et le contrôle non destructif (CND), nous recherchons un(e) Ingénieur(e) en Intelligence Artificielle à Aix en Provence (13)
1/ Vos Missions
Au sein d'une équipe qui développe des algorithmes IA issus de méthodes de CND afin de garantir un haut niveau de sûreté et de performance d'installation, vous aurez les missions suivantes :
Déployer des méthod…

https://www.linkedin.com/jobs/view/4388329891

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4388329891 · captured: 2026-04-18 · fit_reason: low fit

ALTERNANCE - Ingénieur IA appliquée au Reverse Engineering - F/H

55% fit

Thales · Rennes, Brittany, France (Hybrid) · On-site · linkedin_jobs

ROLE

Lieu : Rennes, France

Construisons ensemble un avenir de confiance

Thales est un leader mondial des hautes technologies spécialisé dans trois secteurs d'activité : Défense & Sécurité, Aéronautique & Spatial, et Cyber & Digital. Il développe des produits et solutions qui contribuent à un monde plus sûr, plus respectueux de l'environnement et plus inclusif. Le Groupe investit près de 4 milliards d'euros par an en Recherche & Développement, notamment dans des domaines clés de l'innovation tels que l'IA, la cybersécurité, le quantique, les technologies du cloud et la 6G. Thales compte près de 81 000 collaborateurs dans 68 pays.

Nos engagements, vos avantages

Notre savoir-faire technologiqueNotre attention portée à l'équilibre des collaborateursUn environnement inclusif et bienveillantUn engagement sociétal et environnemental reconnu (Thales Solidarity, indice CAC 40 ESG…)

Votre quotidien

Vous intégrez l'équipe THALIUM, basée à Rennes, en charge des activités de Lutte Informatique, de recherche de vulnérabilités et de développement d'outils de type Red Team.

Vous contribuez plus spécifiquement aux travaux de l'équipe de rétro-ingénierie en explorant les apports de l'IA au service des experts.

Dans ce contexte, vos missions sont les suivantes :

Réaliser un état de l'art des solutions d'IA appliquées à la rétro-ingénierie (LLM, agents, outils). Assurer une suivi continu sur ces solutions dans un écosystème technologique en évolution rapide. Proposer des contributions en ma…

https://www.linkedin.com/jobs/view/4390882244

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4390882244 · captured: 2026-04-18 · fit_reason: low fit

AI Engineer H/F

Inetum · St.-Ouen, Île-de-France, France (On-site) · On-site · linkedin_jobs

55% fit

ROLE

Description De L'entreprise

Inetum est un leader européen des services numériques. Pour les entreprises, les acteurs publics et la société dans son ensemble, les 28 000 consultants et spécialistes du groupe visent chaque jour l'impact digital : des solutions qui contribuent à la performance, à l'innovation et au bien commun.

Présent dans 19 pays au plus près des territoires, et avec ses grands partenaires éditeurs de logiciels, Inetum répond aux enjeux de la transformation digitale avec proximité et flexibilité.

Porté par son ambition de croissance et d'industrialisation, Inetum a généré en 2023 un chiffre d'affaires de 2,5 milliards d'€.

Pour répondre à un marché en croissance continue depuis plus de 30ans, Inetum a fait le choix délibéré de se recentrer sur 4 métiers afin de gagner en puissance et proposer des solutions sur mesure, adaptées aux besoins spécifiques de ses clients : le conseil (Inetum Consulting), la gestion des infrastructures et applications à façon (Inetum Technologies), l'implémentation de progiciels (Inetum Solutions) et sa propre activité d'éditeur de logiciels (Inetum Software). Inetum a conclu des partenariats stratégiques avec 4 grands éditeurs mondiaux - Salesforce, ServiceNow, Microsoft et SAP et poursuit une stratégie d'acquisitions dédiée afin d'entrer dans le top 5 européen sur ces technologies et proposer la meilleure expertise à ses clients.

Tous nos postes sont ouverts aux personnes en situation de handicap.

Description Du Poste

Concev…

https://www.linkedin.com/jobs/view/4392591832

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4392591832 · captured: 2026-04-18 · fit_reason: low fit

AI Engineer

BEYOND SKILLS · Paris, Île-de-France, France (On-site) · On-site · linkedin_jobs

55% fit

ROLE

Ø=Þ€ AI Engineer : Help Build the Next AI Powerhouse
Location: Paris (Hybrid – 2 days remote/week)Contract: Full-time / PermanentLanguage Requirements: Fluent French + Professional English
Join a fast-growing French AI scale-up, backed by top-tier investors, as an AI Engineer and play a pivotal role in building the technical foundation of the company's product and infrastructure.
This is more than just a technical role it's a unique opportunity to build the architecture, shape the product, and define the culture of a company poised to become a major player in European AI.
Ø>Ýà About the Company
The company has reached a key inflection point: strong product-market fit, growing customer base, and a clear roadmap toward international scaling.
Team today: 15 peopleTarget by year-end: 25+Structured into 3 squads:Software EngineeringArtificial Intelligence (R&D and applied)Product & Design
You'll be at the heart of product development, working hand-in-hand with both technical and business leaders.
Ø=Þà Tech Stack
You'll join a high-performance engineering environment working with modern, scalable technologies:
Frontend: React, TypeScriptBackend: Node.js, NestJSAdditional language: Python, Pytorch (for AI integration & tooling)Database: PostgreSQLAPIs: GraphQL via Hasura
We're building a modular, API-first platform designed to scale your architectural decisions will directly influence its evolution.
Ø=Ý What We're Looking For
We're seeking an experienced engineer who thrives in early-stage…

https://www.linkedin.com/jobs/view/4393877698

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4393877698 · captured: 2026-04-18 · fit_reason: low fit

Gen AI Engineer

WOLSOME · Île-de-France, France (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

Gen AI Engineer (H/F)
Dans le cadre du développement de notre practice Data & AI, nous recherchons un Gen AI Engineer pour intervenir chez nos clients grands comptes sur des projets d'intelligence artificielle générative à fort impact métier (Finance, Assurance, Industrie…).L'objectif de la mission est de concevoir et déployer des applications Gen AI robustes et scalables, en s'appuyant sur les dernières avancées en matière de LLM, d'agents autonomes et de RAG.

MISSIONS :
Conception & développement Gen AI• Analyser les besoins métiers et identifier les cas d'usage pertinents pour la Gen AI (automatisation, assistant, extraction d'information, génération de contenu…).• Concevoir et développer des pipelines RAG (Retrieval-Augmented Generation) : chunking, embedding, vector stores, reranking.• Développer des agents autonomes et des workflows multi-agents avec LangGraph et LangChain.• Intégrer des LLMs (OpenAI, Anthropic, Mistral, Llama…) via API ou en déploiement on-premise selon les contraintes client.
Agentic AI & orchestration• Architecturer des systèmes agentiques capables de raisonner, planifier et agir de façon autonome (ReAct, Tool Use, Chain-of-Thought).• Implémenter des mémoires conversationnelles, des outils externes et des boucles de feedback dans les agents.• Optimiser les prompts et évaluer la qualité des sorties LLM (RAGAS, TruLens, évaluations maison).• Assurer la mise en production des applications Gen AI sur AWS (Bedrock, Lambda, ECS).
Collaboration & documenta…

https://www.linkedin.com/jobs/view/4400552717

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4400552717 · captured: 2026-04-18 · fit_reason: low fit

AI Research & Prototyping Engineer

55% fit

Hornetsecurity France · Greater Paris Metropolitan Region (Hybrid) · On-site · linkedin_jobs

ROLE

Hornetsecurity is a leading European cloud based cybersecurity provider, protecting more than 125,000 organizations in over 120 countries with innovative security, compliance and backup solutions. Founded in 2007 and headquartered in Hannover, Germany, the company has grown and generated over $50 million in revenue while maintaining rapid international expansion and technological leadership in cloud security.
We are looking for an AI Research Engineer to join the PAIR team and play a central role in the design, development, experimentation, and production deployment of innovative AI-driven services (LLMs, autonomous agents, automation, NLP, analytics services, etc.).
You will be involved in the entire project lifecycle, from ideation and prototyping to production deployment, while actively contributing to the applied research initiatives of the Security Lab. You will work closely with multidisciplinary teams to turn ambitious ideas into concrete, robust, and scalable AI solutions in an environment that encourages continuous innovation.
Your mission
End-to-End Ownership of AI SolutionsDesign, develop, and maintain AI services from prototype to production.Ensure robustness, performance, scalability, and operational reliability of solutions in industrial settings.Rigorous Experimentation & Applied ResearchMethodically test and benchmark AI models (standards, metrics, comparisons).Document results and propose innovative solutions tailored to cybersecurity challenges.Innovation & …

https://www.linkedin.com/jobs/view/4377119886

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4377119886 · captured: 2026-04-18 · fit_reason: low fit

Cloud / DevOps Engineer (AI Infrastructure)

55% fit

AI4I · Greater London, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

This opportunity is located in Turin, Italy

The Italian Institute of Artificial Intelligence (AI4I) is seeking a senior, hands-on Cloud / DevOps Engineer to design, build, and operate the cloud foundation of its HPC/AI infrastructure.You will own the Cloud layer that enables AI and HPC workloads to run reliably at scale, with a strong focus on OpenStack-based private cloud infrastructure and the evolution toward container-based environments such as Kubernetes, Rancher, and Harvester.This role is central to AI4I's industrial deployment mission: you will create and operate the infrastructure that powers real AI projects for enterprises, public institutions, and strategic partners.
Location: AI4I, OGR – Turin, ItalyHybrid work: Flexible arrangements may be negotiatedThe position will remain open until filled and multiple candidates may be hired.
About The Role
As Cloud / DevOps Engineer at AI4I, you will take ownership of the end-to-end cloud infrastructure, from architecture and automation to day-to-day operations and continuous improvement.Beyond compute orchestration, this role includes responsibility for the design and operational management of distributed storage systems that support AI and HPC workloads. You will ensure that storage performance, durability, and scalability meet the needs of GPU-intensive training, fine-tuning, and inference environments.This is a cross-unit role supporting multiple AI4I teams and projects, working closely with engineering, deploymen…

https://www.linkedin.com/jobs/view/4399811762

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399811762 · captured: 2026-04-18 · fit_reason: low fit

Cloud Platform Security Engineer

55% fit

CommonAI CIC · Cambridge, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

CommonAI CIC is a non-profit membership organisation, founded on a belief in collaborative engineering for the safe and responsible development of foundational AI technologies. A place where AI startups, enterprises large and small, public sector bodies and academia can share resources and knowledge, to codevelop and grow businesses, fast.

You will be responsible for ensuring that our platform is scalable, secure and resilient. We are looking for a security expert who can be responsible for security across our IT, software development and production environments.

We work with multiple cloud vendors so need someone with deep knowledge of platform security who can apply it across familiar and unfamiliar cloud environments. Specific responsibilities will include:

Working with the CommonAI software team to ensure our environment is secure by designImplementing security best practices across the company for all IT, development and production environmentsWorking with customers to ensure that their use of our platform and services is secure

Requirements

A practical, hands on approach and ability to work across layers of the stack when neededStrong experience with cloud infrastructure and platform security on AWS and other hyperscalersDeep understanding of IAM & identity federationNetworking expertise (VPC/VNet, subnets, routing, VPNs, private endpoints)Knowledge of compute orchestration (Kubernetes, VMs, serverless)Cloud cost optimisation and resource governanceProgramming expe…

https://www.linkedin.com/jobs/view/4389769389

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4389769389 · captured: 2026-04-18 · fit_reason: low fit

Azure Cloud Infrastructure Engineer

55% fit

Royal Mail · Sheffield, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

Job reference 325328

Azure Cloud Infrastructure Engineer

Competitive salary plus car allowance, healthcare , 18% annual bonus potential, 25 days annual leave plus bank holidays rising with service and a company pension scheme with highly competitive contribution rates

Permanent, Full time

Royal Mail Sheffield Hybrid working x3 days in Sheffield office

With a long and proud history, Royal Mail is one of the most visible and trusted brands in the UK. We keep people and businesses connected every hour of every day and we’re passionate about harnessing technology to deliver the best possible results for our customers.

Our business is thriving by continuously adapting and adopting advanced technology. We reach every delivery point in the UK and have over 30 million customer touch points per day. We have 165,000 employees, 80,000 PDA’s and 30,000 end IT users. We are in an exciting period of transformation with significant growth in parcel & packet volumes and online services and we’re focusing on investing in the security, technology, processes, and people that are going to help us continue to achieve great things and grow our business.

We have the UK’s largest “Feet on the Street” network of over 85,000 postmen and women and the lowest reported carbon emissions per parcel of any major UK delivery company. We are committed to providing sustainable services as we deliver against our environmental ambitions. We are investing in technologies such as electric vehicles and drone…

https://www.linkedin.com/jobs/view/4399793108

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399793108 · captured: 2026-04-18 · fit_reason: low fit

Cloud Security Engineer

55% fit

Flo Health Inc. · London, England, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

500M+ downloads. 80M+ monthly users. A decade of building – and we’re still accelerating.

Flo is the world’s #1 health & fitness app worldwide on a mission to build a better future for female health. Backed by a $200M investment led by General Atlantic, we became the first product of our kind to reach a $1B valuation in 2024 – and we’re not slowing down.

With 7M paid subscribers and the highest-rated experience in the App Store’s health category, we’ve spent 10 years earning trust at scale. Now, we’re building the next generation of digital health – AI-powered, privacy-first, clinically backed – to help our users know their body better.

The job

The Scale of the Challenge

At Flo we don't just have users, we have a global community. We are the #1 women's health app, in the last month alone, we saw 8.6M new installs and a 2.8M increase in active users.

When millions of people trust you with their most personal health data, security isn't a feature — it's a foundation. We are looking for a Cloud Security Engineer to join Velocity, our Internal Platform team. Your mission is to ensure that every system, pipeline, and tool our engineers rely on is secure by default — so they can ship fast without ever compromising trust.

The Mission: Velocity

The Velocity team exists to eliminate friction. We build and own the foundation everything else runs on: cloud infrastructure, developer tooling, and SRE practices. You will:

Embed Security Into the Platform: Bake security, compliance…

https://www.linkedin.com/jobs/view/4389264608

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4389264608 · captured: 2026-04-18 · fit_reason: low fit

Infrastructure Engineer (Linux) - £110k-£130k + bonus

55% fit

Vertus Partners · London Area, United Kingdom (Hybrid) · On-site · linkedin_jobs

ROLE

IT Infrastructure Engineer – Buy-Side Technology – London (Hybrid)
We are partnering with a buy-side firm in London that is looking to hire an Infrastructure Engineer to help build, improve, and take ownership of their internal infrastructure environment. This is a hands-on engineering role with a strong focus on Linux, automation, and infrastructure improvement rather than pure support.
The firm is in the process of bringing more infrastructure capability in-house and reducing reliance on external providers, so this role offers a genuine opportunity to influence how infrastructure is managed, automated, and developed going forward.
The OpportunityYou will be responsible for maintaining and improving the infrastructure that supports trading systems, development environments, and internal platforms. The role sits alongside a platform team working on cloud and platform tooling, and you will work closely with them while focusing more on core infrastructure, Linux systems, storage, and data centre environments.This role would suit someone who enjoys fixing inefficiencies, automating manual processes, and introducing better ways of managing infrastructure rather than just keeping systems running.
What You'll Be DoingThe role will involve a mix of infrastructure engineering, automation, and project work. This includes managing Linux servers and storage environments, improving monitoring and resilience, automating operational tasks, and working on infrastructure projects such as sto…

https://www.linkedin.com/jobs/view/4391410750

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4391410750 · captured: 2026-04-18 · fit_reason: low fit

Infrastructure Engineer

SISL Global · London, England, United Kingdom (On-site) · On-site · linkedin_jobs

55% fit

ROLE

Position OverviewSeeking an Infrastructure Engineer with 9-10 years of experience to build and configure GCP services for customer monitoring platforms. Must have expertise in GCP, Terraform, automation, monitoring, and FinOps.
Key ResponsibilitiesDesign, build, and deploy GCP infrastructure for customer monitoring solutionsDevelop and maintain Terraform scripts for infrastructure provisioning and automationImplement monitoring solutions using Cloud Monitoring, Prometheus, Grafana, or DatadogAutomate infrastructure operations using Python, Bash, CI/CD pipelinesConduct FinOps activities: cost optimization, committed use discounts, budgetingConfigure GCP services: Compute Engine, GKE, Cloud Storage, Cloud SQL, VPC, IAM, Cloud FunctionsImplement security best practices, IAM policies, and compliance standardsBuild cost monitoring dashboards and implement labeling strategiesCreate technical documentation, runbooks, and architecture diagramsTroubleshoot infrastructure incidents and optimize performance
Required SkillsExperience:9-10 years in cloud/infrastructure engineering5+ years GCP experience (mandatory)3+ years Terraform experience (mandatory)

Technical Must-Haves:GCP: Compute Engine, GKE, Cloud Storage, Cloud SQL, VPC, IAM, Cloud Functions, Cloud MonitoringTerraform: Modules, state management, IaC automationScripting: Python, Bash, PowerShellCI/CD: Cloud Build, Jenkins, GitHub Actions, GitLab CIMonitoring: Cloud Monitoring, Cloud Logging, Prometheus, Grafana, ELK StackContai…

https://www.linkedin.com/jobs/view/4401063179

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4401063179 · captured: 2026-04-18 · fit_reason: low fit

Google Cloud Engineer

55% fit

Marks Sattin · Glasgow, Scotland, United Kingdom (On-site) · On-site · linkedin_jobs

ROLE

We are seeking a skilled Cloud Engineer with deep Google Cloud Platform (GCP), Python, and Terraform experience to support a large-scale GCP enablement programme. You will design and deliver secure, scalable, automated cloud foundations, enable Vertex AI services, provision production and delivery environments, and help implement the target-state operating model across people, process and technology. Success in this role directly impacts enterprise readiness for AI/ML workloads and platform resilience.
Key responsibilities (decisions & impact-focused)Design, build and automate GCP infrastructure using Terraform and Python to deliver secure, production-ready environments.Provision and manage delivery and Production environments for programme workloads.Create reusable Terraform modules, templates and engineering patterns to accelerate scalable GCP adoption.Implement platform capabilities to enable Vertex AI APIs and enterprise AI/ML workloads.Build automation for provisioning, configuration validation, compliance checks, operational workflows and platform integrations.Translate control-plane and operating-model requirements into deployable engineering assets, standards and CI/CD pipelines.Collaborate with architecture, security, network, IAM and platform teams to deliver compliant, resilient solutions for regulated workloads.Drive environment setup, service onboarding, policy alignment and release readiness for GCP workloads.Document platform designs, implementation patterns an…

https://www.linkedin.com/jobs/view/4399961403

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399961403 · captured: 2026-04-18 · fit_reason: low fit

Senior Platform Reliability Engineer (Operations)

55% fit

Ricoh Europe · London, England, United Kingdom (On-site) · On-site · linkedin_jobs

ROLE

Change your job, change your workplace, change your future...

Ricoh are currently recruiting for a Senior Platform Reliability Engineer (Operations) based in London. This role is a senior, hands on engineering role responsible for the reliability, resilience, and operational integrity of Ricoh's hybrid and cloud platforms.

The role focuses on operating, stabilising, and continuously improving live platforms, ensuring services meet defined availability, performance, security, and compliance standards. It combines deep operational expertise with strong automation, Infrastructure as Code (IaC), and observability capability to reduce toil, improve recovery, and enable predictable service outcomes.

Ricoh transforms organisations, using innovative technologies and services enabling you as an individual to work smarter. This is what we call "empowering digital workplaces".

In fact the entire Ricoh workforce enjoys our pioneering and innovative ways of working. We like to call it: imagine. change., it's the ethos of our brand and how we drive positive change for ourselves and others. Our teams are embracing change, fostering new ways of working and we have never been more resolute in our mission - "you work for us, and we`ll work for you".

What you will be doing

Deliver standards for availability, latency, performance, capacity, and scalability.Take part in root-cause analysis and problem management for major incidents.Champion blameless post-mortem culture and ensure actions a…

https://www.linkedin.com/jobs/view/4389201646

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4389201646 · captured: 2026-04-18 · fit_reason: low fit

AI Operations Developer

Docebo · Toronto, ON (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

Artificial Intelligence. Actual Impact.

At Docebo, we're using AI to change how people learn at work—and we mean actually change it. We're an AI-powered learning platform that helps organizations create, deliver, and manage training all in one place. But our real mission goes deeper: we help teams move faster, work smarter, and focus on the work that truly matters. Our platform is built with intelligent, time-saving tools that personalize learning, eliminate busywork, and turn training from a checkbox into a superpower. The result? Better experiences for learners and real results for businesses.

We're shaping the future of learning with a team that isn't afraid to challenge the status quo. If you're excited by the idea of using AI to make work-life better for real people–you'll feel right at home here. And it's not just what we build, it's how we show up. At Docebo, our values aren't just posters on the wall—they guide how we work every day. We call it the Docebo Heart: trust by default, assume positive intent, and create space for different perspectives to thrive.

So… what are you waiting for? Join 900+ Docebians around the world and help us reinvent the way people learn, because learning never stops.

The Adventure Ahead Ø=Þ€

You are joining a small, high-impact team at the forefront of the AI revolution where we don't just talk about the future—we build it. As an AI Operations Developer, your mission is to transform Docebo's internal engines by shipping AI-native applica…

https://www.linkedin.com/jobs/view/4399483137

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399483137 · captured: 2026-04-18 · fit_reason: low fit

AI Automation Developer

55% fit

Great Foods - Egyptian International For Food Products · Montreal, QC (Hybrid) · On-site · linkedin_jobs

ROLE

Company DescriptionGreat Foods is one of Egypt's largest food companies and a key part of the El Naggar Group. Founded in 1978, Great Foods employs over 2000 individuals across six factories. We specialize in producing a wide variety of products including meats, fruits, jams, juices, and canned goods, marketed under the brand names Al Marai, OK, Juicy Up, and Royal. Our range of over 100 products caters to both local and international markets.
Great Foods is going all in on AI.We're building internal tools that automate the way we work — and we're not looking for just any coder. We want someone who can use AI to build code (not sit and write it manually all day). Think GitHub Copilot, Claude, and other generative AI tools as your partners in crime.This is not a back-office dev role — you'll be front and center, walking through every department, understanding how people actually get work done, and automating the repetitive parts of their job.You'll speak to logistics, QA, sales, accounting — learn their workflows — and then create smart, AI-powered tools to remove inefficiencies. We're giving you the room to build solutions that scale with our growing company.This is a full-time hybrid role based in Montreal. You'll be expected to work 3–4 days per week in the office, alongside the team.Ø=Þàþ What You'll DoShadow and map out workflows across every departmentInterview team members to understand their daily responsibilities and pain pointsUse AI tools (ChatGPT, Copilot, etc.) to g…

https://www.linkedin.com/jobs/view/4397477844

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4397477844 · captured: 2026-04-18 · fit_reason: low fit

Operational AI Automation Developer

55% fit

Softchoice · Toronto, ON (On-site) · On-site · linkedin_jobs

ROLE

Why You'll Love Softchoice

We are a software-focused IT solutions and services provider that equips organizations to be agile and innovative, and for their people to be engaged, connected, and creative at work. That means moving them to the cloud, helping them build the workplace of tomorrow, and enabling them to make smarter decisions about their technology. By doing these things we help them create success for their customers and their people.

We stand proudly for our people and support their success through career development and advancement. We are recognized and respected for our culture of inclusion and belonging, continuously striving to do what's good for our people and communities.

The Impact You Will Have

The Operational AI Automation Builder is a hands-on AI builder who accelerates operational modernization by turning solution designs into shipped scalable automations across the Microsoft Power Platform. Reporting to the Director, Customer Success & Operations Improvement, you will partner closely with a Sr. Operational Automation & AI Specialist (solution architect/process designer) to execute step-by-step automation designs using Power Apps, Power Automate, Copilot Studio, and Microsoft's broader AI & automation toolset.

This role is ideal for an advanced citizen developer who thrives on delivery: building dependable workflows, apps, and copilots that reduce manual effort, improve quality, and create repeatable operational capabilities. You will also contrib…

https://www.linkedin.com/jobs/view/4390990991

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4390990991 · captured: 2026-04-18 · fit_reason: low fit

AI Engineer (RAG & LLM Systems)

Wissen Technology · Montreal, QC (On-site) · On-site · linkedin_jobs

55% fit

ROLE

About the job
The Wissen team continues to expand its footprint in the USA, Canada, UK, Australia, and India. More openings to come as we continue to grow the team!
Please read below for a brilliant career opportunity.
Role: AI Engineer (RAG & LLM Systems)Title: AVP/VPLocation: Montreal, QC(Day 1 Onsite/Hybrid) – FulltimeMode of Work: 3 days/week onsite requiredRequired Experience: 5-10+ years
We are hiring an AI Engineer with strong hands-on experience in RAG systems, LLM APIs, and AI agents. This role focuses on building production-ready AI applications, including retrieval systems and intelligent automation workflows.
Required Skills:AG / Retrieval Systems - Hands-on experience building RAG pipelines. Familiarity with NotebookLM, RAG tooling, and Python libraries (NotebookLM or similar).LLM APIs & Tooling -Experience with Claude / Anthropic APIs, Claude Code, or equivalent LLM tooling. Comfortable building production systems on top of LLM APIs.Vector Databases -Exposure to Pinecone, Chroma, or similar. Understanding of embedding-based retrieval architecture and indexing strategies.AI Orchestration- Lang Chain or similar frameworks for multi-step AI workflows. Chaining, tool use, and prompt management at scale.AI Agents & Automation-Building AI agents or workflow automations beyond simple chat interfaces. Multi-step reasoning, tool integration, autonomous task execution.LLM Governance-Awareness of governance - logging, monitoring, cost controls, and responsible AI practices…

https://www.linkedin.com/jobs/view/4393599552

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4393599552 · captured: 2026-04-18 · fit_reason: low fit

AI Engineer

Scotiabank · Toronto, ON (On-site) · On-site · linkedin_jobs

55% fit

ROLE

Requisition ID: 255027

Join a purpose driven winning team, committed to results, in an inclusive and high-performing culture.

We are seeking a passionate and execution-focused AI Engineer to join our high-impact engineering team in Global Wealth Management Technology (GWMT). This role is ideal for an individual contributor who thrives in a fast-paced, agile environment and is ready to design, build, and deploy production-grade AI solutions with minimal supervision. You will be surrounded by innovators solving some of the most pressing challenges in wealth management through intelligence-driven systems.

Scotiabank's Global Wealth Management (GWM) division plays a critical role in delivering personalized financial advice, investment management, and estate planning solutions to clients across Canada and around the world. As part of this trusted and forward-looking organization, you'll join a high-performing technology team focused on building intelligent, secure, scalable platforms that enhance client and advisor experiences.

Working within GWM means applying your AI expertise to real-world financial challenges—helping shape the next generation of wealth services through data-driven innovation and responsible engineering. If you're motivated by meaningful impact,

collaboration, and continuous learning, this is where your work will matter.

Is this role right for you? In this role you will:

Design, implement, and manage Generative AI use cases for production use.Engineer ro…

https://www.linkedin.com/jobs/view/4387126673

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4387126673 · captured: 2026-04-18 · fit_reason: low fit

Agentic AI Engineer

Arcadia · Toronto, ON (On-site) · On-site · linkedin_jobs

55% fit

ROLE

A well-funded, fast-growing AI company is seeking an Agentic AI Engineer to design and deploy intelligent, agent-based systems that tackle complex, multi-step workflows. The team is building next-generation AI tools capable of reasoning, planning, and executing tasks autonomously across sophisticated technical domains.In this role, you will develop advanced agent architectures that integrate retrieval, orchestration, and reasoning, turning traditionally manual or fragmented processes into seamless, end-to-end AI-driven solutions. This position offers a unique opportunity to work at the intersection of research and production, with direct impact on cutting-edge AI capabilities.
What You'll DoDevelop multi-step reasoning agents incorporating memory, retrieval-augmented generation (RAG), and external tool/API integrationsDesign and implement evaluation frameworks to measure and improve agent performance, including reasoning accuracy and system reliabilityBuild scalable orchestration patterns, including supervisor and multi-agent coordination strategiesDeploy production-grade agentic systems in cloud environments (AWS or similar)Partner with cross-functional teams to translate complex workflows into robust, agent-driven solutions
What We're Looking ForExperience building agent-based AI systems, LLM-powered workflows, or orchestration pipelinesStrong understanding of retrieval systems, reasoning techniques, RAG, and multi-agent architecturesHands-on experience deploying AI solutio…

https://www.linkedin.com/jobs/view/4395333875

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4395333875 · captured: 2026-04-18 · fit_reason: low fit

Director, Agentic AI Solutions Architect

Mastercard · Toronto, ON (Hybrid) · On-site · linkedin_jobs

55% fit

ROLE

Our Purpose

Mastercard powers economies and empowers people in 200+ countries and territories worldwide. Together with our customers, we’re helping build a sustainable economy where everyone can prosper. We support a wide range of digital payments choices, making transactions secure, simple, smart and accessible. Our technology and innovation, partnerships and networks combine to deliver a unique set of products and services that help people, businesses and governments realize their greatest potential.

Title And Summary

Director, Agentic AI Solutions Architect

Overview

Services within Mastercard is responsible for acquiring, engaging, and retaining customers by managing fraud and risk, enhancing cybersecurity, and improving the digital payments experience. We deliver value added services by combining traditional management consulting with Mastercard’s data, AI assets, proprietary platforms, and execution expertise to provide clients with powerful strategic insights and recommendations.

The Mastercard Agent Suite team is building what’s next at the intersection of AI, commerce, and payments. Launched in 2026, Agent Suite is a global product that enables enterprises to build, test, and deploy customizable agentic AI systems within their own infrastructure. As agentic AI moves from experimentation to production, Mastercard occupies a unique position of trust, network reach, and advisory depth to make this transition safe, responsible, and impactful for the world’s leading …

https://www.linkedin.com/jobs/view/4397954706

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4397954706 · captured: 2026-04-18 · fit_reason: low fit

Applied AI, Forward Deployed Machine Learning Engineer - Montreal

55% fit

Mistral AI · Montreal, QC (Hybrid) · On-site · linkedin_jobs

ROLE

About Mistral

At Mistral AI, we believe in the power of AI to simplify tasks, save time, and enhance learning and creativity. Our technology is designed to integrate seamlessly into daily working life.

We democratize AI through high-performance, optimized, open-source and cutting-edge models, products and solutions. Our comprehensive AI platform is designed to meet enterprise needs, whether on-premises or in cloud environments. Our offerings include le Chat, the AI assistant for life and work.

We are a dynamic, collaborative team passionate about AI and its potential to transform society.

Our diverse workforce thrives in competitive environments and is committed to driving innovation. Our teams are distributed between France, USA, UK, Germany and Singapore. We are creative, low-ego and team-spirited.

Join us to be part of a pioneering company shaping the future of AI. Together, we can make a meaningful impact. See more about our culture on https://mistral.ai/careers.

About The Job

Mistral AI is seeking a Applied AI Engineer to facilitate the adoption of its products among customers and collaborate with them to address complex technical challenges.

The Applied AI Engineer will be an integral part of our Applied AI Engineering team, which is dedicated to driving the successful deployment of Mistral AI products. They will work hand-in-hand with customers from the pre-sale stage to post-implementation, ensuring our solutions meet and exceed client expectations.

In this r…

https://www.linkedin.com/jobs/view/4388756138

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4388756138 · captured: 2026-04-18 · fit_reason: low fit

AI Enablement Engineer

55% fit

Electronic Arts (EA) · Vancouver, BC (Hybrid) · On-site · linkedin_jobs

ROLE

Description & Requirements

Electronic Arts creates next-level entertainment experiences that inspire players and fans around the world. Here, everyone is part of the story. Part of a community that connects across the globe. A place where creativity thrives, new perspectives are invited, and ideas matter. A team where everyone makes play happen.

Our EA Experiences group (XO) is dedicated to ensuring great experiences for our growing communities centered around our world-renowned brands, including fan-favorites like Apex, Battlefield, EA SPORTS FC, Madden NFL and The Sims, just to name a few. We're a multi-functional group, with world-class expertise building fandoms, driving interactive storytelling, and positioning our franchises at the center of the broader entertainment ecosystem. We inspire, connect, and engage fans through culturally relevant content, intentionally architected journeys across channels, and meaningful fan care. Our goal is to provide valuable, easy experiences that fans love – in our games, around our games, and through innovative adjacent experiences to grow and enrich how fans experience EA as we shape the future of entertainment.

To empower more players and fans in new and amazing ways, we need more innovators to join our world-class team. The future of entertainment is interactive, and you can help lead that future, by growing and enriching how hundreds of millions of people (and counting) find joy and belonging, forge friendships, and celebrate th…

https://www.linkedin.com/jobs/view/4344151524

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4344151524 · captured: 2026-04-18 · fit_reason: low fit

AI Builder / Forward Deployed Engineer

55% fit

Sagard · Toronto, ON (On-site) · On-site · linkedin_jobs

ROLE

Sagard Overview

Sagard is a multi-strategy alternative asset management firm active in venture capital, private equity, private credit and real estate. Sagard also engages in private wealth management through Sagard Wealth.

Founded in 2016 and guided by the core values of entrepreneurship, innovation, collaboration, rigour and authenticity, Sagard has experienced outstanding growth. Today, the firm has more than US$45 billion under management, 190 portfolio companies and 540 professionals.

Sagard is well positioned to continue to grow substantially, organically and inorganically, pursuing its vision of becoming one of the best-performing investment management firms. Headquartered in Canada, Sagard currently has offices in Canada, the United States, Europe and the Middle East.

More at https://www.sagard.com

Position Overview

Sagard is on a journey to become an AI-first organization. We believe artificial intelligence will reshape the workplace by automating repetitive tasks and enabling teams to focus on higher-value, strategic work.

The AI Builder / Forward Deployed Engineer will work closely with business teams across investment, finance, compliance and operations to identify high-value opportunities and build practical AI-enabled solutions. As part of a small and agile AI and Data team, the role combines technical development, experimentation and cross-functional collaboration to rapidly prototype, deploy and scale solutions that improve how work is done across the f…

https://www.linkedin.com/jobs/view/4382460771

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4382460771 · captured: 2026-04-18 · fit_reason: low fit

Founding Engineer

50% fit

Nozomio · San Francisco, CA, US · On-site · ycombinator

Ø=Ü° $150K - $250K

ROLE

Founding Engineer at Nozomio | Y Combinator Nozomio index anything. search it with ai. Founding Engineer $150K - $250K • 2.00% - 4.00% • San Francisco, CA, US Job type Full-time Role Engineering, Full stack Experience Any (new grads ok) Visa US citizen/visa only Skills Python, TypeScript, Google Cloud Connect directly with founders of the best YC-funded startups. Apply to role › Arlan Rakhmetzhanov Founder Arlan Rakhmetzhanov Founder About the role About Nia and Nozomio Nozomio is applied product and research lab. Our flagship product Nia is an API layer that provides agents with up-to-date, continuously monitored context from libraries, research papers, and technical documentation—saving you hours of manual ingestion. Today, agents are very good at generating code but still bad at knowing which repo, package, or doc to look at. Context, not generation, is the bottleneck. Nia exists to solve that bottleneck so agents reach near superhuman code intelligence. Nozomio is backed by Y Combinator and recently raised a 6.2M seed round from CRV, BoxGroup, LocalGlobe, YC, 20VC, and angels such as Paul Graham, Thomas Wolf, Gustav Söderström, and Gokul Rajaram. We are building a very small, very strong team in San Francisco. The goal is to hire a small number of exceptional people before IPO, not scale to thousands. You will likely get a job if... Can spend $1,000 on Cursor or Claude Code in 1–2 days (my record is $3,000 for production-ready features). Have a green GitHub (no bs contrib…

https://www.ycombinator.com/companies/nozomio/jobs/blaVouI-founding-engineer

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:82620 · captured: 2026-04-18 · fit_reason: low fit

Founding Engineer

50% fit

Replicas · San Francisco, CA, US · On-site · ycombinator

Ø=Ü° $150K - $250K

ROLE

Founding Engineer at Replicas | Y Combinator Replicas Background coding agents for engineering teams Founding Engineer $150K - $250K • 2.00% - 4.00% • San Francisco, CA, US Job type Full-time Role Engineering, Full stack Experience Any (new grads ok) Visa Will sponsor Connect directly with founders of the best YC-funded startups. Apply to role › Connor Loi Founder Connor Loi Founder About the role Replicas lets you manage background coding agents - swarms of frontier coding agents like Claude Code or Codex working in their own sandboxes. With their own customizable environments, they can iterate and work on any engineering task and deliver PRs. Replicas comes from the vision proposed by AI 2027 , whereby so much of humanity's day-to-day work will be done by these fleets of AI agents, each equipped with the tools to complete tasks for anything from deep research to software engineering. Our mission is to build the infrastructure and tooling to control these background agents. We want to build a team of exceptional people in San Francisco, to work on some of the most interesting problems and products of the next decade! About You We're looking for people who chase as hard as they can. We want someone with opinions about where AI agents and software engineering is heading. Someone who has so many ideas, can build every step of the way, talk to customers, and go deep on technical problems. You are a fit if: You have startup experience (e.g founder / founding eng) You have a very,…

https://www.ycombinator.com/companies/replicas/jobs/1r7KcWj-founding-engineer

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92797 · captured: 2026-04-18 · fit_reason: low fit

Engineering Manager, MVR

50% fit

Checkr · San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $232K - $274K

ROLE

Engineering Manager, MVR at Checkr | Y Combinator Checkr People infrastructure for the future of work Engineering Manager, MVR $232K - $274K • San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the Role The MVR team builds the systems behind driver record checks — one of Checkr's highest-volume screening products, relied on by transportation, gig economy, and logistics companies to make fast, confident hiring decisions for millions of candidates each year. The portfolio spans standard MVR, Continuous MVR (ongoing driver monitoring), Commercial MVR, and DOT compliance tools including CDLIS and the Drug & Alcohol Clearinghouse, covering all 51 DMV jurisdictions across the US, Puerto Rico, and Canada. As Engineering Ma…

https://www.ycombinator.com/companies/checkr/jobs/UzLoBEc-engineering-manager-mvr

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92268 · captured: 2026-04-18 · fit_reason: low fit

Software Engineer, Reliability

50% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $127K - $176K

ROLE

Software Engineer, Reliability at Checkr | Y Combinator Checkr People infrastructure for the future of work Software Engineer, Reliability $127K - $176K • Denver, Colorado, United States; San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 3+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . As a Software Engineer II on the Site Reliability Engineering team within the Platform Engineering group at Checkr, you will identify reliability challenges impacting engineering teams and platforms and develop innovative solutions to resolve them. You will strive towards the right balance between enforcing standardization and accommodating tailored workflows. The person in this role will have the opportunity to demonstrate a degree of autonomy in their work and help with complex support…

https://www.ycombinator.com/companies/checkr/jobs/RhiskBr-software-engineer-reliability

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92296 · captured: 2026-04-18 · fit_reason: low fit

Engineering Manager, Truework

50% fit

Checkr · San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $232K - $274K

ROLE

Engineering Manager, Truework at Checkr | Y Combinator Checkr People infrastructure for the future of work Engineering Manager, Truework $232K - $274K • San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 11+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the role The Truework Team at Checkr is looking for a motivated Engineering Manager to lead our Mortgage Verifications Team. This team is responsible for building and launching new products, customer-facing experiences, and enhancing the Truework platform. As the Engineering Manager for this team, you will work hand-in-hand with product managers and key stakeholders to create the software and experiences that delight our customers. The projects you'll be working on are high impact and scale: you'll be advancing our…

https://www.ycombinator.com/companies/checkr/jobs/7ibHUDQ-engineering-manager-truework

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92269 · captured: 2026-04-18 · fit_reason: low fit

Staff Product Manager, Identity Risk Solutions

50% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $176K - $244K

ROLE

Staff Product Manager, Identity Risk Solutions at Checkr | Y Combinator Checkr People infrastructure for the future of work Staff Product Manager, Identity Risk Solutions $176K - $244K • Denver, Colorado, United States; San Francisco, California, United States Job type Full-time Role Product Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . In an AI-driven world, distinguishing real from fake is increasingly challenging. AI has made it easy to forge credentials, impersonate individuals, and submit job applications at scale using synthetic identities. More than a third of remote applicants are fake, and Fortune 500 companies experienced a significant rise in hires from state-sponsored fraudulent actors last year. AI has turned the hiring funnel from a protective filter into an open gateway for bad actors. Che…

https://www.ycombinator.com/companies/checkr/jobs/nnkIzOj-staff-product-manager-identity-risk-solutions

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92306 · captured: 2026-04-18 · fit_reason: low fit

Staff Product Manager, Growth

50% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $176K - $244K

ROLE

Staff Product Manager, Growth at Checkr | Y Combinator Checkr People infrastructure for the future of work Staff Product Manager, Growth $176K - $244K • Denver, Colorado, United States; San Francisco, California, United States Job type Full-time Role Product Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . In this role as a Staff Product Manager, Growth, you will lead the Growth team and deliver products and capabilities for Checkr's growth funnel, with a focus on signups, conversion, user activation and lifetime value per customer for the company's self-serve business. As the Staff Product Manager for Growth, you will play a pivotal role in driving the expansion of our user base while ensuring an exceptional user experience throughout the entire customer journey. Roles and responsibilities include creating…

https://www.ycombinator.com/companies/checkr/jobs/3cO9by6-staff-product-manager-growth

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92305 · captured: 2026-04-18 · fit_reason: low fit

Sr. Growth Marketing Manager

50% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $133K - $186K

ROLE

Sr. Growth Marketing Manager at Checkr | Y Combinator Checkr People infrastructure for the future of work Sr. Growth Marketing Manager $133K - $186K • Denver, Colorado, United States; San Francisco, California, United States Job type Full-time Role Marketing Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . About the team/role Checkr is hiring a Senior Growth Marketing Manager to drive pipeline and self-serve revenue growth for our Truework business. This role will report to Checkr's Sr. Director of Growth Marketing and will be instrumental in shaping and scaling acquisition, onboarding, and retention through web-led acquisition, product-led activation, and lifecycle engagement. This role will be the bridge between Marketing and Product, driving growth through experimentation across the entire user journey. …

https://www.ycombinator.com/companies/checkr/jobs/VF48jxA-sr-growth-marketing-manager

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92297 · captured: 2026-04-18 · fit_reason: low fit

Senior Software Engineer, Trust

50% fit

Checkr · San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $185K - $218K

ROLE

Senior Software Engineer, Trust at Checkr | Y Combinator Checkr People infrastructure for the future of work Senior Software Engineer, Trust $185K - $218K • San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . Checkr Trust is building safer communities by providing our customers with risk signals about its users. Every year, our team provides insights on tens of millions of users across multiple verticals: online dating, car rental, short-term rentals, and many more. We acquire, analyze, and deliver insights in smart, scalable ways that help organizations identify and mitigate risk. Trust is part of the Checkr start-up lab: we operate with the speed of a small startup, while enjoying the leverage and resources of a known te…

https://www.ycombinator.com/companies/checkr/jobs/bqjjBE5-senior-software-engineer-trust

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92289 · captured: 2026-04-18 · fit_reason: low fit

Backend Software Engineer II

50% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $127K - $176K

ROLE

Backend Software Engineer II at Checkr | Y Combinator Checkr People infrastructure for the future of work Backend Software Engineer II $127K - $176K • Denver, Colorado, United States; San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 3+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . As a Software Engineer at Checkr, you will work on high impact engineering projects that help build a fairer future for all. You will work in a collection of small services built on Ruby and Javascript, with both SQL and NoSQL databases, as well as message queues for asynchronous tasks. Checkr was founded by engineers for engineers; we want to create the best environment for developers to build, learn and love their work. What you'll do: Own the product development of features from design, i…

https://www.ycombinator.com/companies/checkr/jobs/exXko1l-backend-software-engineer-ii

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92264 · captured: 2026-04-18 · fit_reason: low fit

Senior Software Engineer, Reliability

50% fit

Checkr · Denver, Colorado, United States; San Francisco, California, United States · On-site · ycombinator

Ø=Ü° $157K - $218K

ROLE

Senior Software Engineer, Reliability at Checkr | Y Combinator Checkr People infrastructure for the future of work Senior Software Engineer, Reliability $157K - $218K • Denver, Colorado, United States; San Francisco, California, United States Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Checkr and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Checkr Checkr is building the data platform to power safe and fair decisions. Established in 2014, Checkr's innovative technology and robust data platform help customers assess risk and ensure safety and compliance to build trusted workplaces and communities. Checkr has over 100,000 customers including DoorDash, Coinbase, Lyft, Instacart, and Airtable. We're a team that thrives on solving complex problems with innovative solutions that advance our mission. Checkr is recognized on Forbes Cloud 100 2025 List and is a Y Combinator 2024 Breakthrough Company . As a Senior Software Engineer on the Site Reliability Engineering team within the Platform Engineering group at Checkr, you will identify issues and technical challenges, develop innovative solutions to resolve them, and drive their adoption. You will strive towards the right balance between enforcing standardization and accommodating tailored workflows. The person in this role will have the opportunity to demonstrate a degree of autonomy in their work and help with complex…

https://www.ycombinator.com/companies/checkr/jobs/61vR0nx-senior-software-engineer-reliability

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:92290 · captured: 2026-04-18 · fit_reason: low fit

Software Engineer, Computer Vision and Deep Learning

50% fit

Mashgin · Palo Alto, CA · On-site · ycombinator

ROLE

Software Engineer, Computer Vision and Deep Learning at Mashgin | Y Combinator Mashgin Self-Checkout using Computer Vision. Software Engineer, Computer Vision and Deep Learning • Palo Alto, CA Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Mashgin and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Mashgin At Mashgin, we are developing the future of checkout experiences. We aim to reduce the friction between customer desires and their ability to be on their way. Our ecosystem is anchored by the only 3D computer vision and deep learning based checkout system in the world. With customers across a wide spectrum of verticals including the NFL, MLB, and Aramark, our company and product adoption are growing at a rapid pace. We value autonomy of ones work, a culture of respect for one another, and building products with empathy for our customers. Position Summary Mashgin is looking for a passionate Computer Vision engineer who loves solving everyday vision problems. Someone who is regularly saying to themselves “This could be done with computer vision!” and who loves experimenting with new algorithms. You’ll work on a small dynamic team to design and implement cutting edge Computer Vision algorithms that will be used everyday by our customers. You Will Be Developing new computer vision algorithms with founders in C/C++ and Python for solving challenging real-world problems …

https://www.ycombinator.com/companies/mashgin/jobs/Lwg6Tpa-software-engineer-computer-vision-and-deep-learning

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:49135 · captured: 2026-04-18 · fit_reason: low fit

Senior Technical Product Manager

50% fit

Mashgin · Palo Alto, CA, US · On-site · ycombinator

Ø=Ü° $180K - $230K

ROLE

Senior Technical Product Manager at Mashgin | Y Combinator Mashgin Self-Checkout using Computer Vision. Senior Technical Product Manager $180K - $230K • Palo Alto, CA, US Job type Full-time Role Product Experience 6+ years Visa Will sponsor Apply to Mashgin and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Mashgin Mashgin powers the world's best checkout experience for over 40 million users. Customers just place their items on our kiosks and our AI rings up their entire order in less than a second. With Mashgin, lines are now optional. Mashgin's technology powers over 1 billion transactions at your favorite locations, including over half of all US professional sports teams, 4,000 convenience stores, major airports, universities, and more. We're not just building cutting-edge AI—we're creating real-world impact and unforgettable experiences. Backed by a well-funded Series B, we're also one of the rare AI startups that's already profitable. Our secret? A culture of extreme ownership, autonomy, and customer obsession. At Mashgin, you're not just an employee—you're an owner. We might do things differently than most startups… but then again, most startups aren't profitable with an army of raving fans behind them. Position Summary As a Senior Technical Product Manager at Mashgin you will help define products and features that integrate software, hardware, and services to provide frictionless commerce experiences. This role re…

https://www.ycombinator.com/companies/mashgin/jobs/o1C6SdM-senior-technical-product-manager

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:64562 · captured: 2026-04-18 · fit_reason: low fit

Senior Software Engineer, Full-Stack

Mashgin · Palo Alto, CA · On-site · ycombinator

50% fit

ROLE

Senior Software Engineer, Full-Stack at Mashgin | Y Combinator Mashgin Self-Checkout using Computer Vision. Senior Software Engineer, Full-Stack • Palo Alto, CA Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Mashgin and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Mashgin At Mashgin, we are developing the future of checkout experiences. We aim to reduce the friction between customer desires and their ability to be on their way. Our ecosystem is anchored by the only 3D computer vision and deep learning based checkout system in the world. With customers across a wide spectrum of verticals including the NFL, MLB, and Aramark, our company and product adoption are growing at a rapid pace. We value autonomy of ones work, a culture of respect for one another, and building products with empathy for our customers. Position Summary We're looking for a full-stack developer to help us build the software that runs our kiosks and powers the extensive infrastructure behind them. In addition to building more traditional web stacks, you'll work on code that interfaces with our hardware and visual components, making sure kiosks are extra reliable and the user experience amazing. This is not a specialized position -- we believe in end-to-end ownership of projects, so you'll be involved throughout the entire product lifecycle, from idea generation, design, and prototyping to executio…

https://www.ycombinator.com/companies/mashgin/jobs/BaWLmBF-senior-software-engineer-full-stack

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:49134 · captured: 2026-04-18 · fit_reason: low fit

Senior Software Engineer, Backend

50% fit

Mashgin · Palo Alto, CA · On-site · ycombinator

Ø=Ü° $225K - $350K

ROLE

Senior Software Engineer, Backend at Mashgin | Y Combinator Mashgin Self-Checkout using Computer Vision. Senior Software Engineer, Backend $225K - $350K • Palo Alto, CA Job type Full-time Role Engineering, Backend Experience 6+ years Visa US citizen/visa only Apply to Mashgin and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role About Mashgin At Mashgin, we are developing the future of checkout experiences. We aim to reduce the friction between customer desires and their ability to be on their way. In May 2022, Mashgin closed a Series B funding round at a $1.5B valuation. We are backed by some of Silicon Valley's leading investors including Matrix Partners, NEA and Y Combinator. Our ecosystem is anchored by the only 3D computer vision and deep learning based checkout system in the world. We have 4,000+ locations world wide with customers across a wide spectrum of verticals including the NFL, MLB, Circle K, ampm, Aramark and Compass Group. We value autonomy of ones work, a culture of respect for one another, and building products with empathy for our customers. Position Summary We're looking for a backend developer to help us build the software that runs our kiosks and powers the extensive infrastructure behind them. You'll work on code that interfaces with our hardware and visual components, making sure kiosks are reliable, and building scalable backend systems. This is not a specialized position -- we believe in end-to-end owner…

https://www.ycombinator.com/companies/mashgin/jobs/kOoM7gQ-senior-software-engineer-backend

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:67554 · captured: 2026-04-18 · fit_reason: low fit

Customer Success Manager

50% fit

SnapMagic · Redwood City, CA, US / San Francisco, CA, US / Palo Alto, CA, US / San Jose, CA, US / San Mateo, CA, US · On-site · ycombinator

Ø=Ü° $120K - $170K

ROLE

Customer Success Manager at SnapMagic | Y Combinator SnapMagic AI-assisted electronics design Customer Success Manager $120K - $170K • Redwood City, CA, US / San Francisco, CA, US / Palo Alto, CA, US / San Jose, CA, US / San Mateo, CA, US Job type Full-time Role Sales Experience 6+ years Visa US citizen/visa only Connect directly with founders of the best YC-funded startups. Apply to role › Natasha Baker CEO Natasha Baker CEO About the role SnapMagic, the leading platform for electrical engineers, is revolutionizing electronics design with artificial intelligence. We help 1.5 million engineers design products like drones, electric cars, and virtual reality headsets 10x faster. With 10+ years of proven market traction (revenue driven growth), SnapMagic is fresh off a round of VC funding from leading Silicon valley investors to accelerate our business into the next stages of growth. What you'll do: As a pivotal member of our Customer Success team, you'll be at the forefront of driving growth within our customer base at SnapMagic. You will guide global component manufacturers on their digital transformation journey, enabling a new online channel for them to grow their influence within the engineering community. You will be responsible for building value, establishing deep relationships, and driving renewals and expansions within our global customer base of electronic component suppliers. What you'll be doing: Onboarding and Adoption: Guide new customers through the onboarding pr…

https://www.ycombinator.com/companies/snapmagic/jobs/eriIncs-customer-success-manager

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:68675 · captured: 2026-04-18 · fit_reason: low fit

Software Engineer - Deepgram for Restaurants

50% fit

Deepgram · San Francisco, CA · On-site · ycombinator

Ø=Ü° $150K - $230K

ROLE

Software Engineer - Deepgram for Restaurants at Deepgram | Y Combinator Deepgram Building foundational AI for speech transcription and understanding. Software Engineer - Deepgram for Restaurants $150K - $230K • San Francisco, CA Job type Full-time Role Engineering, Full stack Experience 6+ years Visa US citizen/visa only Apply to Deepgram and hundreds of other fast-growing YC startups with a single profile. Apply to role › About the role Company Overview Deepgram is the leading platform underpinning the emerging trillion-dollar Voice AI economy, providing real-time APIs for speech-to-text (STT), text-to-speech (TTS), and building production-grade voice agents at scale. More than 200,000 developers and 1,300+ organizations build voice offerings that are 'Powered by Deepgram', including Twilio, Cloudflare, Sierra, Decagon, Vapi, Daily, Cresta, Granola, and Jack in the Box. Deepgram's voice-native foundation models are accessed through cloud APIs or as self-hosted and on-premises software, with unmatched accuracy, low latency, and cost efficiency. Backed by a recent Series C led by leading global investors and strategic partners, Deepgram has processed over 50,000 years of audio and transcribed more than 1 trillion words. There is no organization in the world that understands voice better than Deepgram. Company Operating Rhythm At Deepgram, we expect an AI-first mindset—AI use and comfort aren't optional, they're core to how we operate, innovate, and measure performance. Every t…

https://www.ycombinator.com/companies/deepgram/jobs/C3qC38V-software-engineer-deepgram-for-restaurants

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:88497 · captured: 2026-04-18 · fit_reason: low fit

Ingénieur.e IA & Business Unit Disruptive

50% fit

CENTURION SEARCH | Cabinet de recrutement · Levallois-Perret, Île-de-France, France (On-site) · On-site · linkedin_jobs

ROLE

Innovation + Esprit d’analyse + Sens business + IA = le combo de votre prochain challenge ?

Rejoignez ce groupe en tant qu' Ingénieur.e IA & Business Unit Disruptive.

Pour une société de conseil en croissance, nous recherchons un.e Ingénieur.e Grande École (X, Mines, Ponts, Centrale, etc.) pour créer et piloter une Business Unit disruptive, dédiée au développement d’activités commerciales au travers d’outils IA à construire et déployer.

En collaboration étroite avec la Direction Générale, en tant que Spécialiste IA, vous serez en charge de concevoir des solutions IA sur mesure, de les industrialiser, puis de les déployer en interne au sein de la future BU que vous serez amené. à piloter.

Ce poste s’adresse à un.e entrepreneur.se dans l’âme, alliant créativité, sens business et pragmatisme, capable de faire émerger des opportunités et de fédérer des équipes autour d’un projet ambitieux.

Conception et développement d’outils IA innovants

Vous analysez les workflows internes (reporting, validations, suivis) pour identifier des gains d’efficacité.Vous concevez, testez et optimisez des prompts complexes pour des cas d’usage internes (RH, commerce, opérations) et externes (clients industriels).Vous développez des agents IA (Copilot, GPT, assistants spécialisés) adaptés aux besoins métiers, en collaboration avec les équipes techniques.Vous prototypiez des outils IA et validez leur pertinence avec la direction avant déploiement.Vous accompagnez les équipes dans la prise en main …

https://www.linkedin.com/jobs/view/4398259708

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4398259708 · captured: 2026-04-18 · fit_reason: low fit

Artificial Intelligence Engineer

50% fit

Meet Life Sciences · Paris, Île-de-France, France (Hybrid) · On-site · linkedin_jobs

ROLE

Role: Tech Lead / Hands-on Lead AI Software Engineer Industry: Life Sciences Location: Paris, Hybrid
We are currently partnered with a team developing advanced AI models that are making personalised treatments a reality at scale. They are looking to hire a hands-on AI/Data Tech Lead to work closely with the CTO and help drive the architecture, delivery, and growth of their platform.
Expectations:Design, build, and scale data-intensive systems and AI/ML-driven platforms handling complex scientific datasetsArchitect and optimise data pipelines, backend systems, and analytical workflowsContribute to the development of a next-generation platform for scientific and biomedical data analysis, ensuring scalability and reliabilityCollaborate closely with data scientists, engineers, and product teams to bring models into productionTake ownership of technical decisions, system design, and engineering best practicesContribute to team leadership, mentoring, and project delivery
Background within the Healthtech/Biomedical/Biotech/Life Sciences Space preffered
Skills:Strong experience with AI/ML systems in production environmentsBackend: Python (FastAPI, Airflow or similar)Architecture: Experience designing scalable, distributed systems and complex data workflowsDatabases & Infrastructure: PostgreSQL, Elasticsearch, Snowflake or similar technologiesCloud & DevOps: AWS/GCP, Docker, Kubernetes, CI/CD pipelinesExperience working with large-scale, data-intensive systemsExposure to event-driven …

https://www.linkedin.com/jobs/view/4391072679

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4391072679 · captured: 2026-04-18 · fit_reason: low fit

Software Engineer, Cloud Infrastructure

OpenAI · London, England, United Kingdom (On-site) · On-site · linkedin_jobs

50% fit

ROLE

About The Team

The Applications Engineering team works across research, engineering, product, and design to bring OpenAI's technology to consumers and businesses.

You'll join the team responsible for running the core infrastructure that supports products like ChatGPT and the API. The systems we support include our kubernetes clusters, infrastructure deployment, our networking stack, cloud abstractions, and more.

We seek to learn from deployment and distribute the benefits of AI, while ensuring that this powerful tool is used responsibly and safely. Safety is more important to us than unfettered growth.

About The Role

The cloud infrastructure team builds and maintains infrastructure abstractions allowing OpenAI to ship products quickly and scalably.

In This Role, You Will

Design and build the development and production platforms that power our products, enabling reliability and security at scaleEnsure our infrastructure can scale to the next order of magnitudeHelp create a diverse, equitable, and inclusive culture that makes all feel welcome while enabling radical candor and the challenging of group thinkLike all other teams, we are responsible for the reliability of the systems we build. This includes an on-call rotation to respond to critical incidents as needed.

You Might Thrive In This Role If You

Have 5+ years building core infrastructureHave experience operating orchestration systems such as Kubernetes at scaleHave experience building abstractions over cloud pla…

https://www.linkedin.com/jobs/view/4306939704

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4306939704 · captured: 2026-04-18 · fit_reason: low fit

Cloud Operation Engineer

50% fit

Centric Software · London, England, United Kingdom (On-site) · On-site · linkedin_jobs

ROLE

About Centric Software

Centric Software® is a global leader, providing an innovative and AI-enabled product-concept-to-commercialization platform for retailers, brands and manufacturers of all sizes. We equip retail, fashion, luxury, footwear, outdoor, home and consumer goods brands with pioneering best-of-breed solutions /to plan, design, develop, source, comply, buy, make, price, allocate, sell and replenish products. Our technology powers brands to streamline processes, drive efficiency and operate with confidence in an ever-changing market.

Our story is one of rapid growth, bold ideas and extraordinary opportunities. We're here to challenge the status quo—and we're looking for brilliant people who want to do the same. No matter where you are in the world, this is your chance to be part of something exceptional.

Job Summary

This role will be responsible for utilizing a diverse set of cloud technologies to support Centric Cloud Customers, and to help build and maintain a scalable, robust, highly available platform for Centric's C8 cloud application.

This position will report to the Senior Director, Cloud Operations.

Responsibilities

Respond to and resolve time-critical customer issues in the public cloudPerform customer deployments, migrations and upgrades in the cloud environmentSupport automation projects including writing codingMonitor, troubleshoot, maintain, and manage servers in multiple cloud environments Create, manage, and maintain Cloud networking components…

https://www.linkedin.com/jobs/view/4369934195

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4369934195 · captured: 2026-04-18 · fit_reason: low fit

Platform Infrastructure Engineer

50% fit

Red Bull Racing & Red Bull Technology · Milton Keynes, England, United Kingdom (On-site) · On-site · linkedin_jobs

ROLE

At Red Bull Technology, peak performance isn't just for the track - it's built behind the scenes. Our Support Functions power every aspect of our success so If you thrive in a fast-paced, high-performance environment and want to be part of something extraordinary, we want to hear from you!
Purpose: Design, build and operate resilient, secure and scalable infrastructure that underpins HPC engineering, DevOps and simulation workloads across on-premises and cloud environments. The role focuses on delivering well-engineered compute, storage, networking and virtualisation foundations using Linux-based systems, automation and Infrastructure-as-Code practices. Working closely with Developer teams, InfoSec and wider IT, this role contributes to the reliable operation and continuous improvement of the infrastructure platforms and services provided. The role also requires a foundational understanding of High Performance Computing (HPC) environments and practical experience supporting datacentre operations, including hardware installation, diagnostics and break/fix activities.
Accountabilities: Design, implement and operate core infrastructure services end-to-end (compute, storage, virtualisation, networking and supporting services) for the Technology CampusBuild and maintain reliable Linux-based infrastructure, including provisioning, patching, hardening and lifecycle managementOwn and continuously improve on-premises and cloud infrastructure environments, proactively monitoring perfor…

https://www.linkedin.com/jobs/view/4399006951

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4399006951 · captured: 2026-04-18 · fit_reason: low fit

Infrastructure Engineer

Experis UK · Cheltenham, England, United Kingdom (On-site) · On-site · linkedin_jobs

50% fit

ROLE

We are partnering with a growing technology organisation to hire an Infrastructure Engineer who can play a key role in designing, delivering, and supporting robust IT infrastructure across both internal and customer environments. This role would suit someone who is hands-on, organised, and eager to broaden their technical skill set. Exposure to networking, automation, virtualisation, or cloud technologies is highly beneficial.
What You'll Be DoingDesigning, building, and supporting physical and virtual infrastructure solutionsInvestigating and resolving infrastructure issues efficientlyManaging system health through backups, patching, monitoring, and lifecycle upgradesProducing clear documentation for infrastructure, processes, and incidentsCreating scripts and automation tools to streamline repetitive tasks
What We're Looking ForDegree in Computer Science/Engineering or equivalent practical experienceBackground in system administration or supportExperience delivering or leading IT infrastructure projectsStrong understanding of networking (Cisco, Juniper, HP, Dell) and core protocols (TCP/IP, DNS, DHCP)Hands-on experience with virtualisation platforms (VMware or Hyper-V)Scripting/IaC skills (Ansible/AWX, Python, PowerShell, Bash)Comfortable with both Windows and Linux administrationKnowledge of backup technologies such as NetBackup or Veeam
Nice to HaveExperience with Git-based source control (GitLab, Bitbucket)Familiarity working within Agile teams
Personal QualitiesHighly o…

https://www.linkedin.com/jobs/view/4324621961

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4324621961 · captured: 2026-04-18 · fit_reason: low fit

AI Engineering Intern, Voice & LLM Systems

20% fit

Presto Phoenix, Inc. · Ontario, Canada (Remote) · Remote ' · linkedin_jobs

ROLE

About Presto Phoenix, Inc.

Presto is the leading Voice AI company for restaurant drive-thrus, operating at scale in complex, noisy, customer-facing environments. As the AI partner to more than a dozen of the most iconic American restaurant brands, Presto is building one of the most impactful real-world applications of AI, directly improving revenue, labor efficiency, and guest experience for millions of people.

Founded by alumni from MIT, the University of Waterloo and University of Illinois, Presto is built on a deeply technical, entrepreneurial engineering culture. We place a premium on ownership, speed, and first principles thinking, empowering engineers to move quickly, experiment boldly, and solve hard, real-world problems.

AI is not a feature at Presto, it is the foundation of how we build, evaluate, and evolve our products. We are backed by Remus Capital, were a Y Combinator company, and are headquartered in Silicon Valley.

This may not be the right fit if you're looking for a passive, low-ownership internship. We move at the pace of AI. Presto is for builders, experimenters, and problem-solvers who thrive in ambiguity and are excited to work on real-world AI systems alongside a high-performance team.

The Role

We are looking for AI Engineering Interns to contribute to the core intelligence behind Presto's Voice AI platform.

This is a hands-on technical internship at the intersection of large language models, real-time systems, and production AI. You will work on…

https://www.linkedin.com/jobs/view/4397660309

DEEP EXTRACTION

HIRING CONTACT

No contact found yet — run scan to pull hiring manager via Netrows.

DRAFT EMAIL

No draft yet — run `node job.mjs draft` to generate.

ACTIONS: **[ ] Send email** [ ] Edit draft [ ] Archive

job_id: netrows_job:4397660309 · captured: 2026-04-18 · fit_reason: low fit